

# Point Processing & Filtering

CS 4243 Computer Vision & Pattern Recognition

Angela Yao

# Recap & Outline

## Last Week

- What is computer vision
- Applications in computer vision
- Why computer vision is difficult
- Historical development of the field
- Digital image capture and representations
- RGB & HSV colour spaces

## Today's Lecture

- Point processing operations: brightness, contrast, normalization, non-linear mapping
- Histogram stretching, equalization, segmentation
- Linear filtering, cross-correlation vs. convolution operations
- Denoising, sharpening, template matching
- Non-linear filtering, e.g. median filters

# Point Processing

Brightness, contrast, normalization, gamma, logarithmic & exponential mapping

Image histograms, histogram stretching & equalization

# Notation

$\mathbf{P}$  : 2D array of pixel data

$p_{ij}$  : element at  $i^{\text{th}}$  of  $I$  rows,  $j^{\text{th}}$  of  $J$  columns

$\mathbf{w}_{ij}$  : window of data centered at  $i, j$

$(0, 0)$   $j$   $J$



# Transformations

Point processing

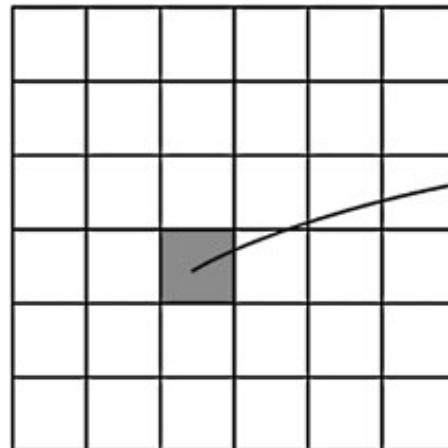
$$x_{ij} = f(p_{ij})$$

corresponding pixel  
from input image

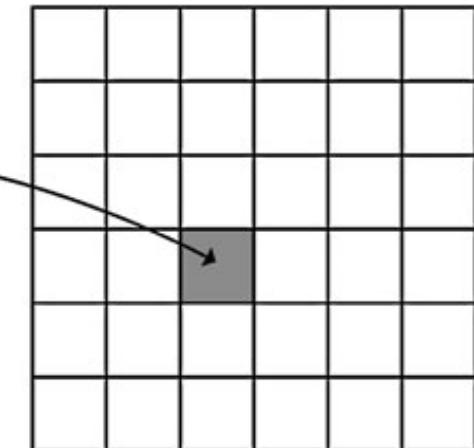
Filtering

$$x_{ij} = f(\mathbf{w}_{ij})$$

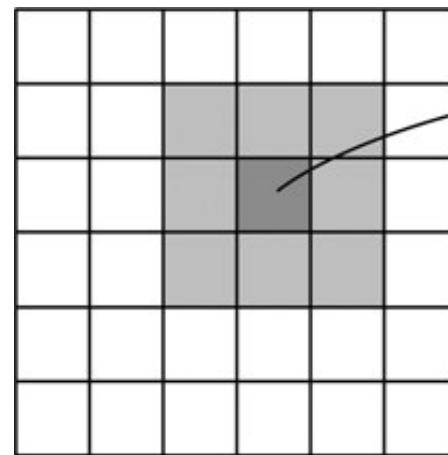
local window of  
data from input



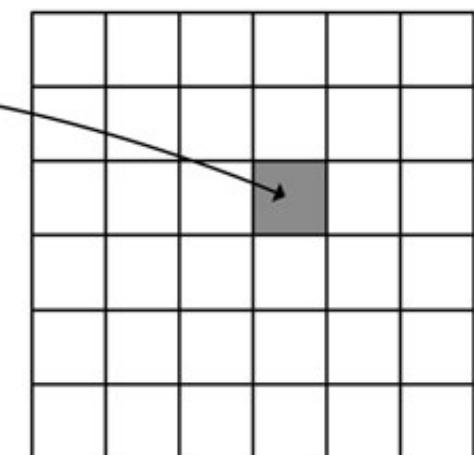
P: input image with elements  $p_{ij}$



X: output image with elements  $x_{ij}$



P: input image with elements  $p_{ij}$



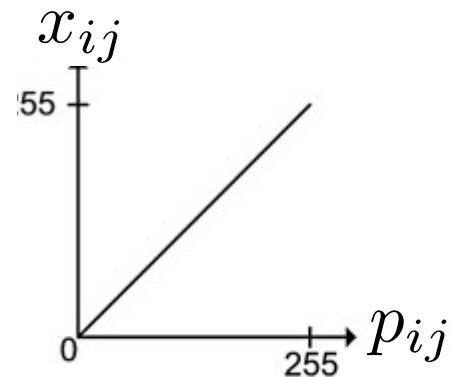
X: output image with elements  $x_{ij}$

# Adjusting Brightness

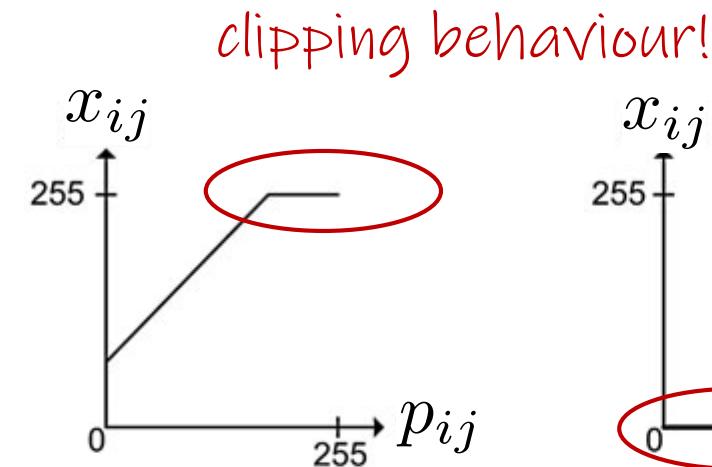
$$x_{ij} = p_{ij} + b$$



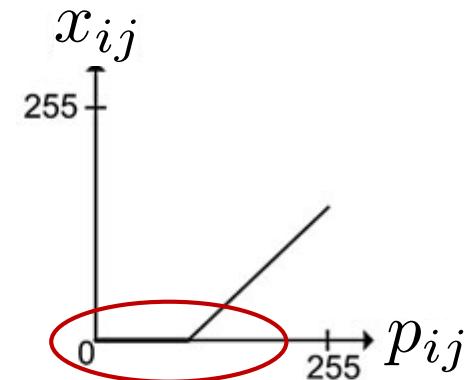
original image



$g(x,y)$ ,  $b = 0$   
no change



$g(x,y)$ ,  $b = 75$   
increase brightness,  $b > 0$



$g(x,y)$ ,  $b = -100$   
decrease brightness,  $b > 0$

# Intensity Scaling

$$x_{ij} = a \cdot p_{ij}$$

$a < 1$



Decreased contrast

$a = 1$



Input image

$a > 1$



Increased contrast

This image on the left is not entirely representative. Scaling with  $a < 1$  decreases the contrast, but this example also has an increase in brightness which is why it looks so grey-ish (otherwise it should look more black). Please have a try with the demo python notebook to get a better idea.

# Image Normalization (Whitening)

$$\text{mean } \mu = \frac{\sum_{i=1}^I \sum_{j=1}^J p_{ij}}{IJ}$$

resulting image pixels are zero-mean, unit variance

$$\text{variance } \sigma^2 = \frac{\sum_{i=1}^I \sum_{j=1}^J (p_{ij} - \mu)^2}{IJ}$$

$$x_{ij} = \frac{p_{ij} - \mu}{\sigma}$$

Removes contrast and constant additive luminance variations

Original

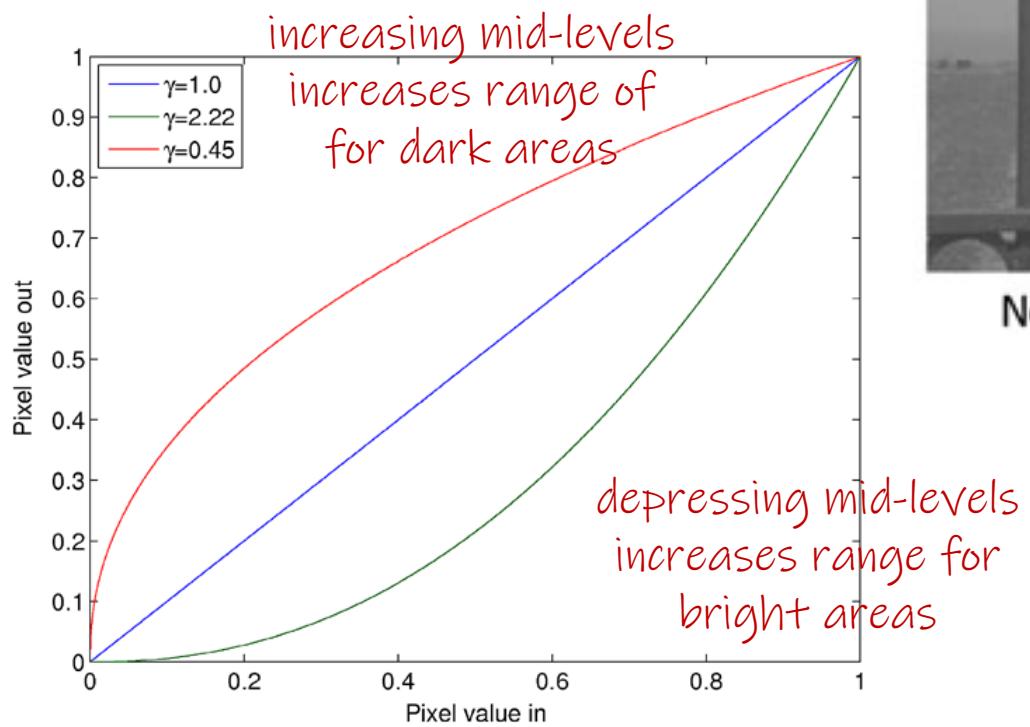


Whitened



# Gamma Mapping

Mappings can also be non-linear, e.g.  $x_{ij} = 255 \cdot \left( \frac{p_{ij}}{255} \right)^\gamma, \gamma > 0$



No gamma correction



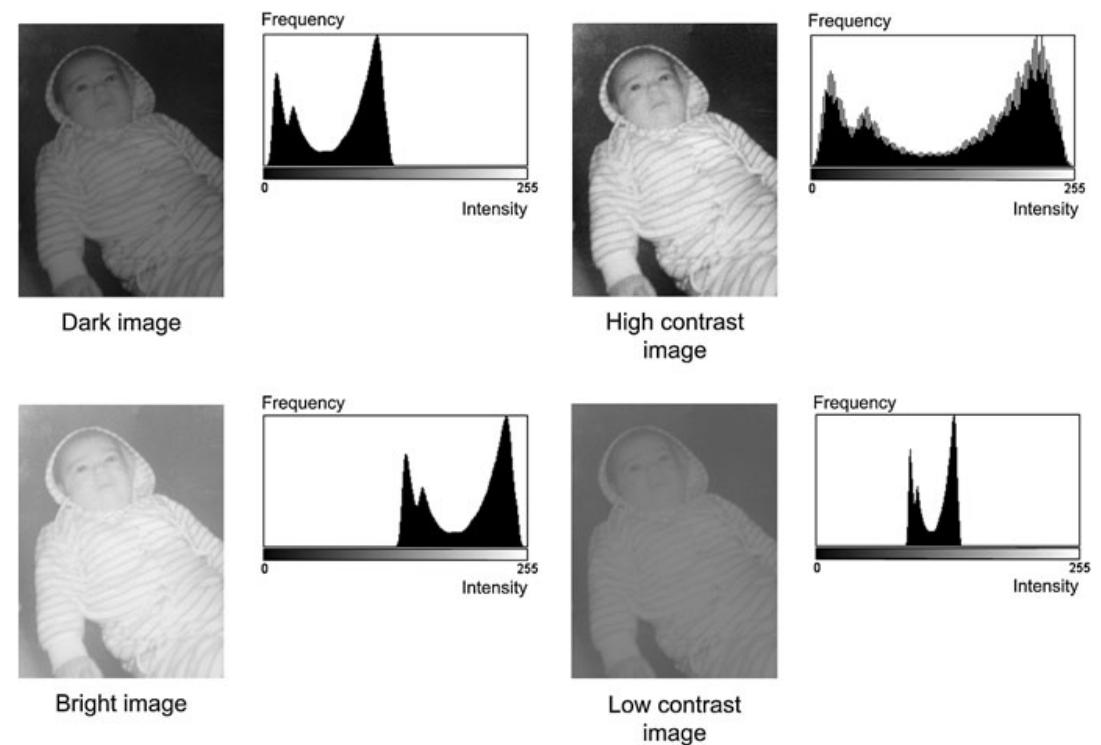
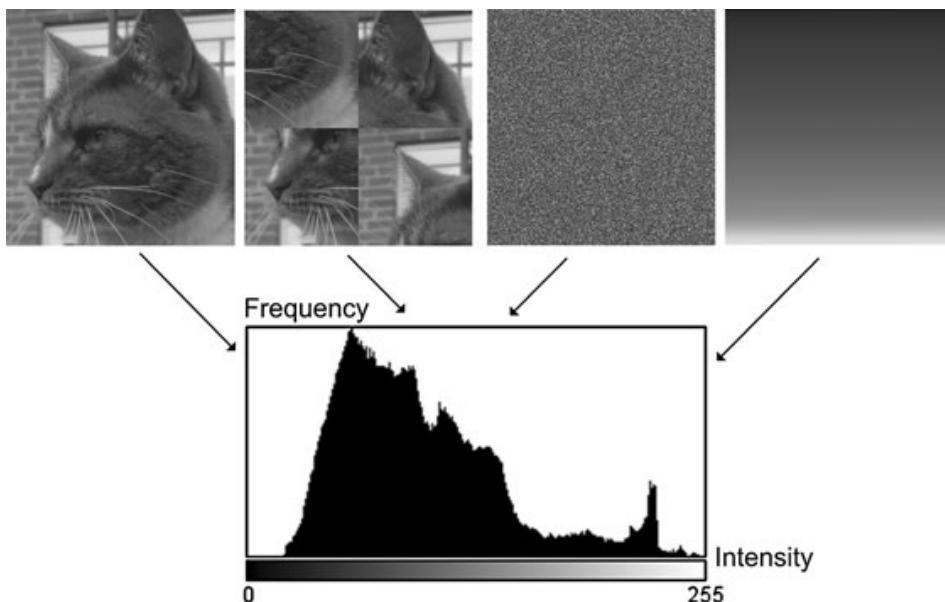
Gamma value: 2.22



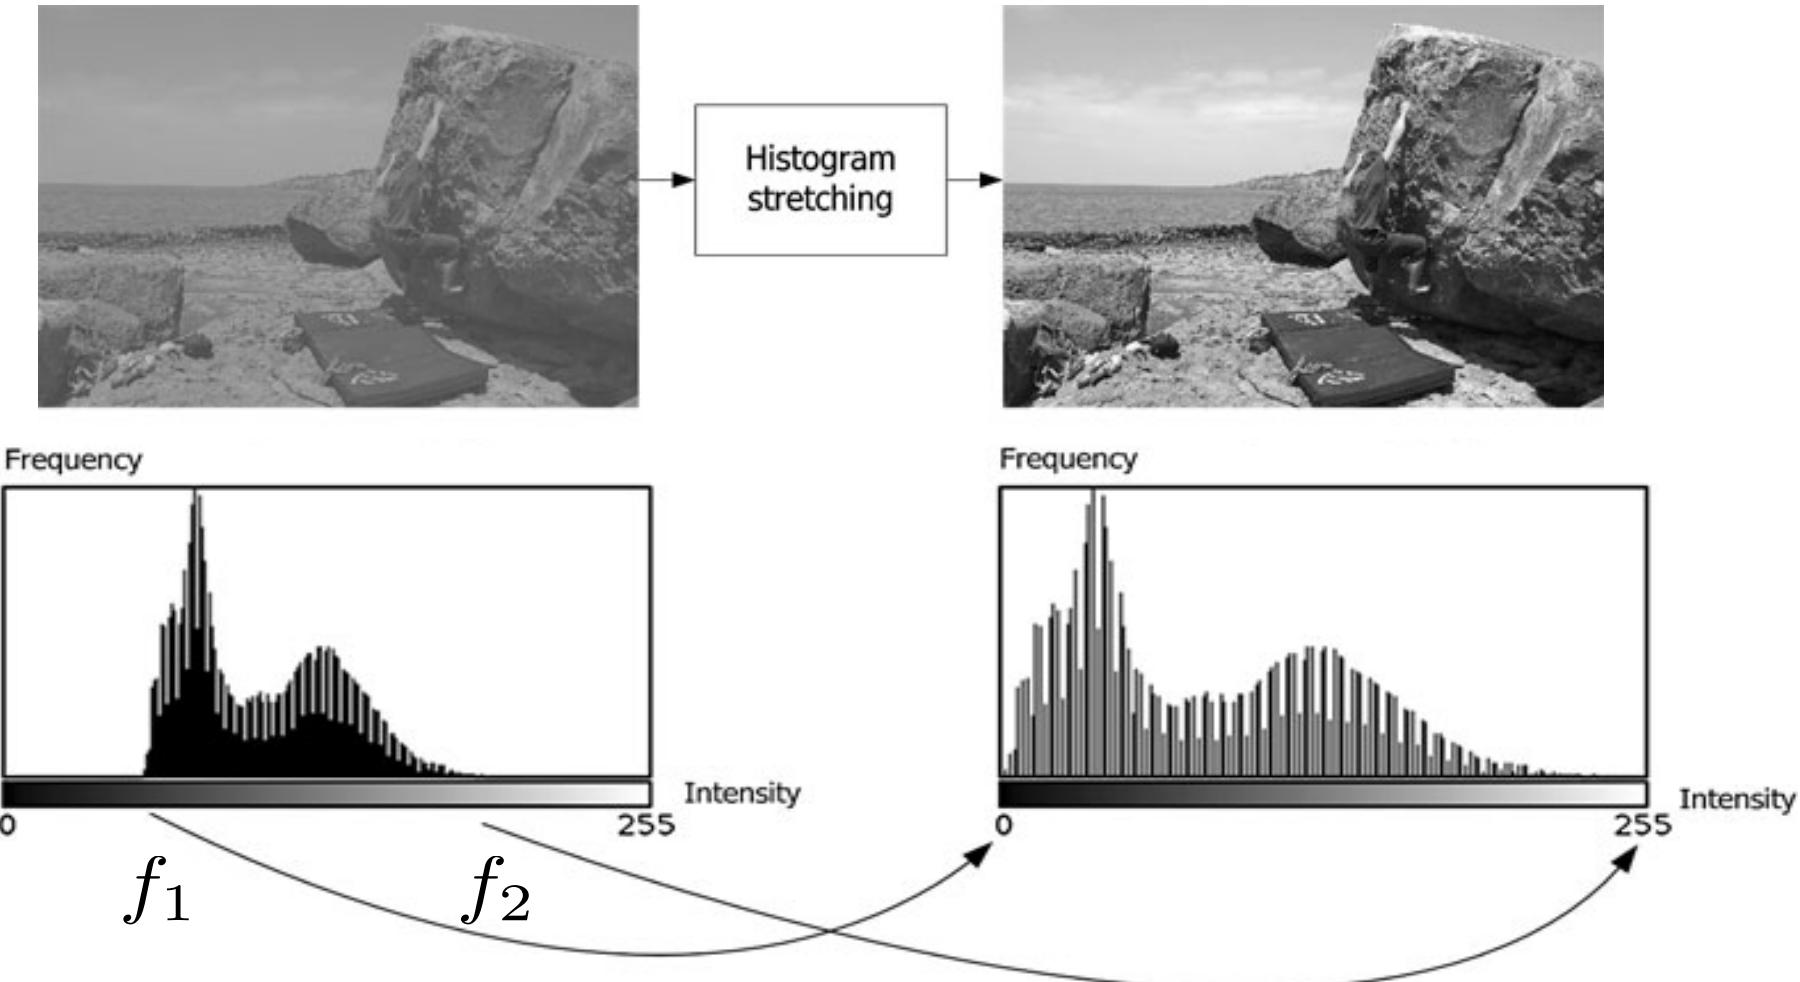
Gamma value: 0.45

# Image Histogram

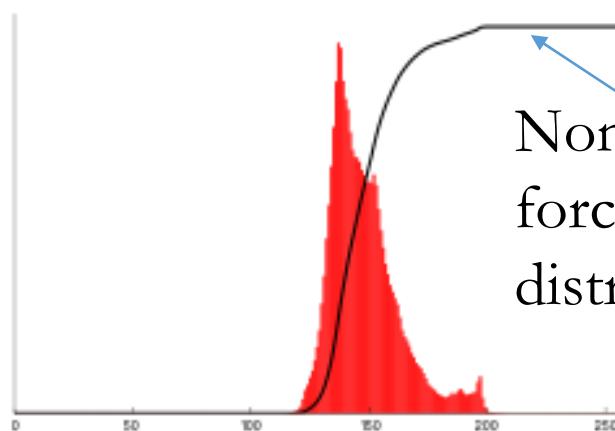
- Information is global (pixel locations don't matter)



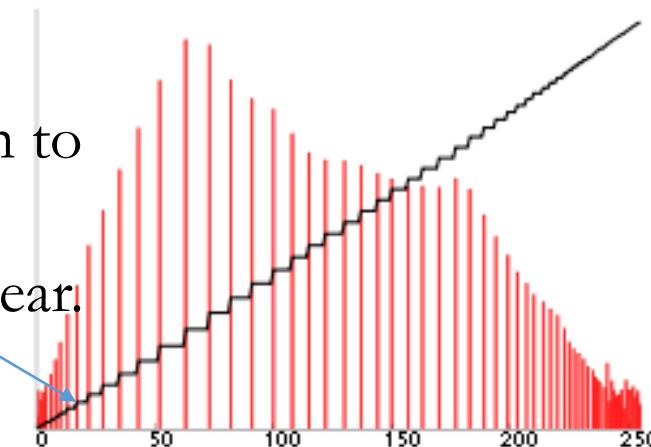
# Histogram Stretching



# Histogram Equalization

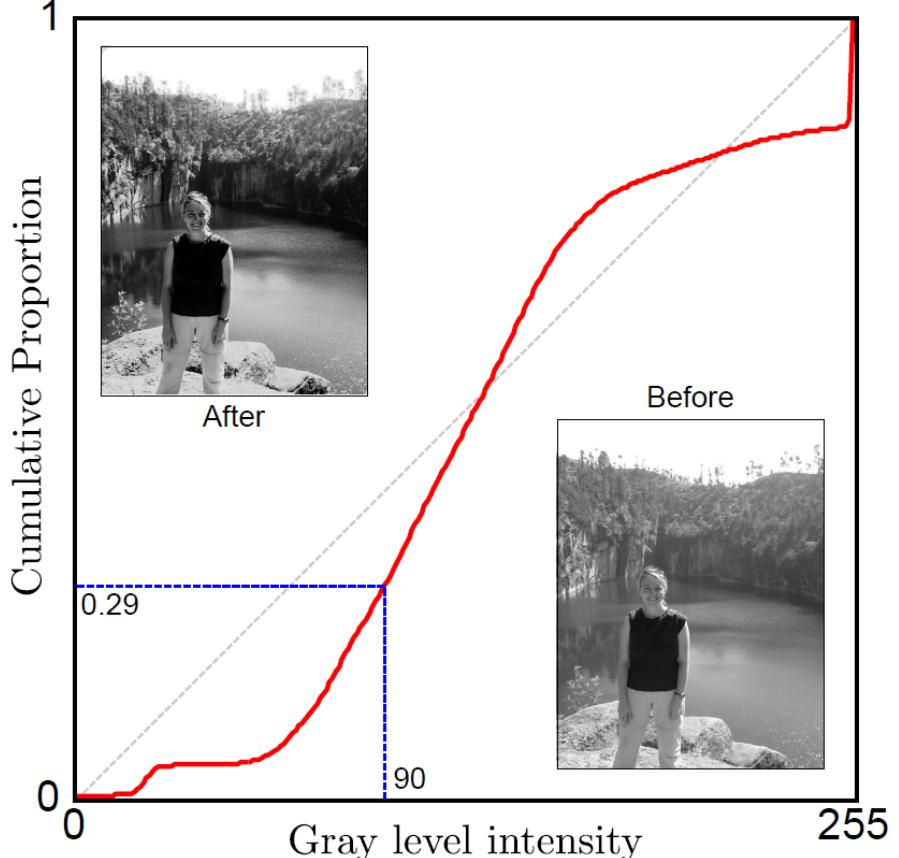


Non-linear operation to force the cumulative distribution to be linear.



# Histogram Equalization

Forces the cumulative distribution to be linear.



histogram

Cumulative sum  
of histogram  
("CDF")

K is usually the  
maximum intensity,  
e.g. 255

Look up for each pixel intensity, the  
corresponding cumulative proportion value  
and remap e.g. 90 maps to  $255 \times 0.29 = 74$

$$h_k = \sum_{i=1}^I \sum_{j=1}^J \delta[p_{ij} - k],$$

$$c_k = \frac{\sum_{l=1}^k h_l}{IJ}.$$

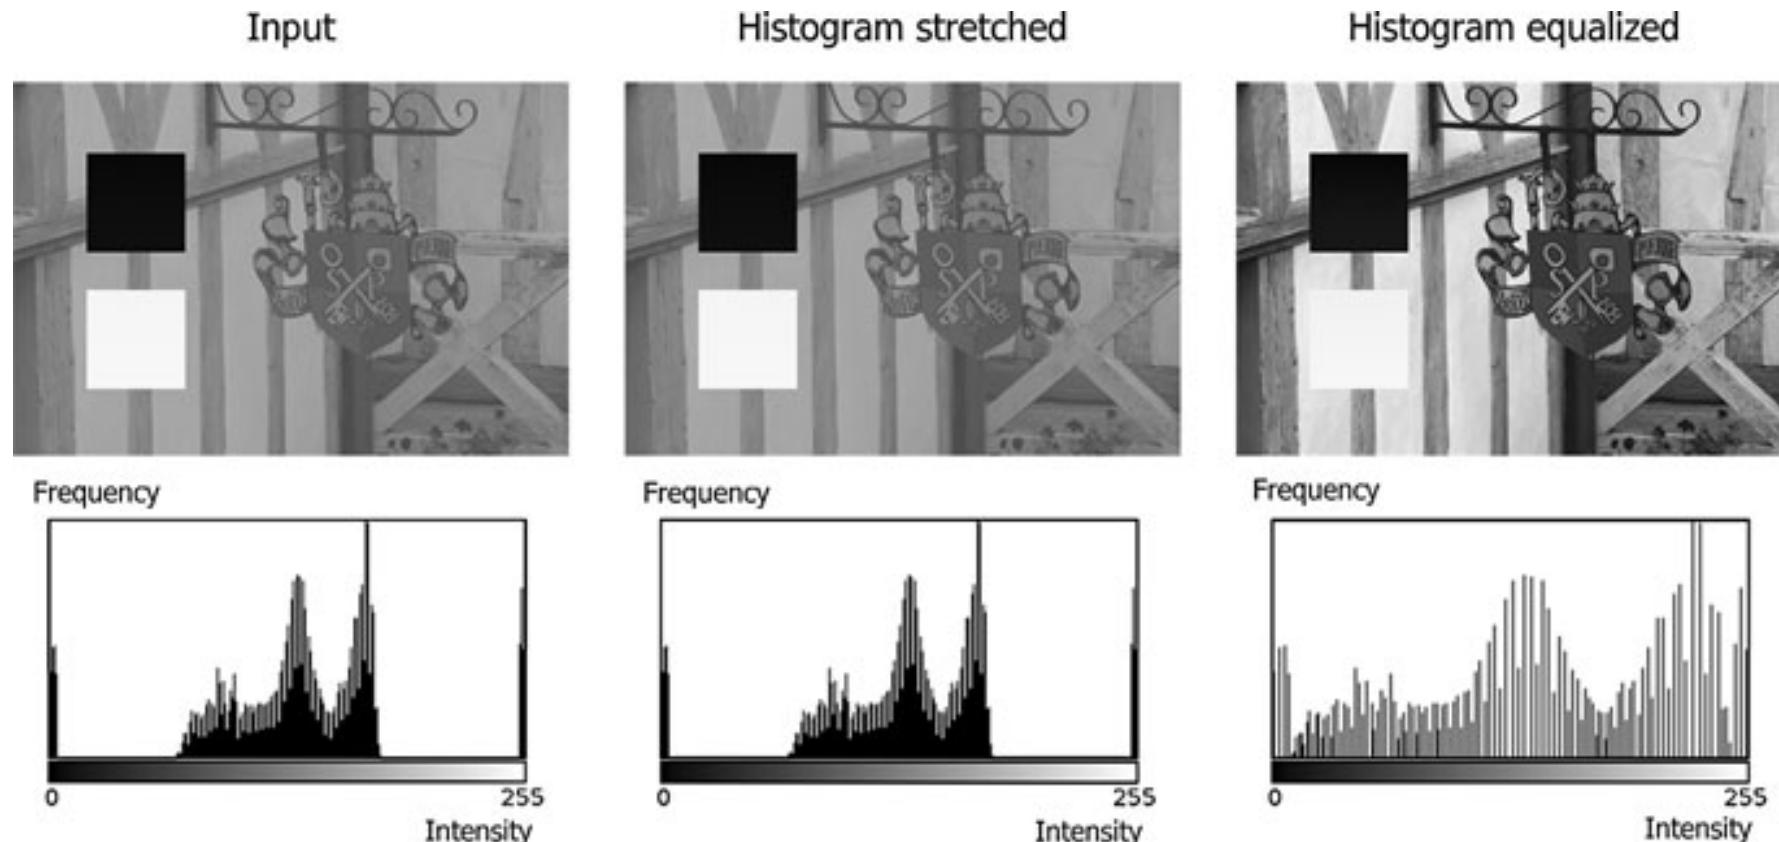
$$x_{ij} = K c_{p_{ij}}.$$

# Equalization vs. Whitening

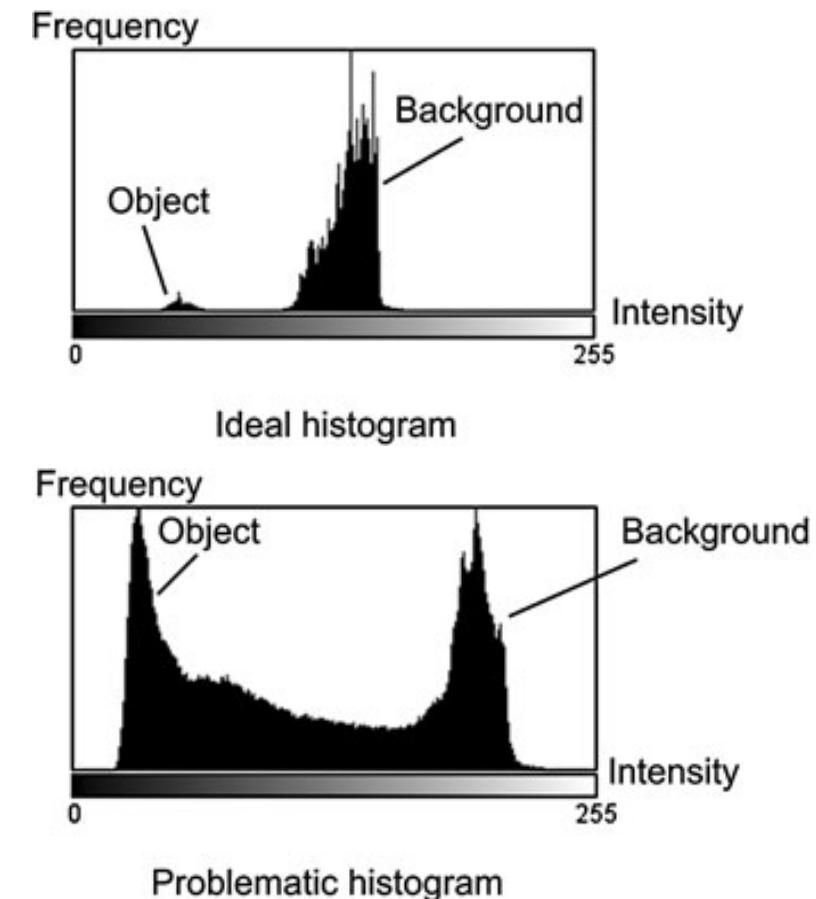
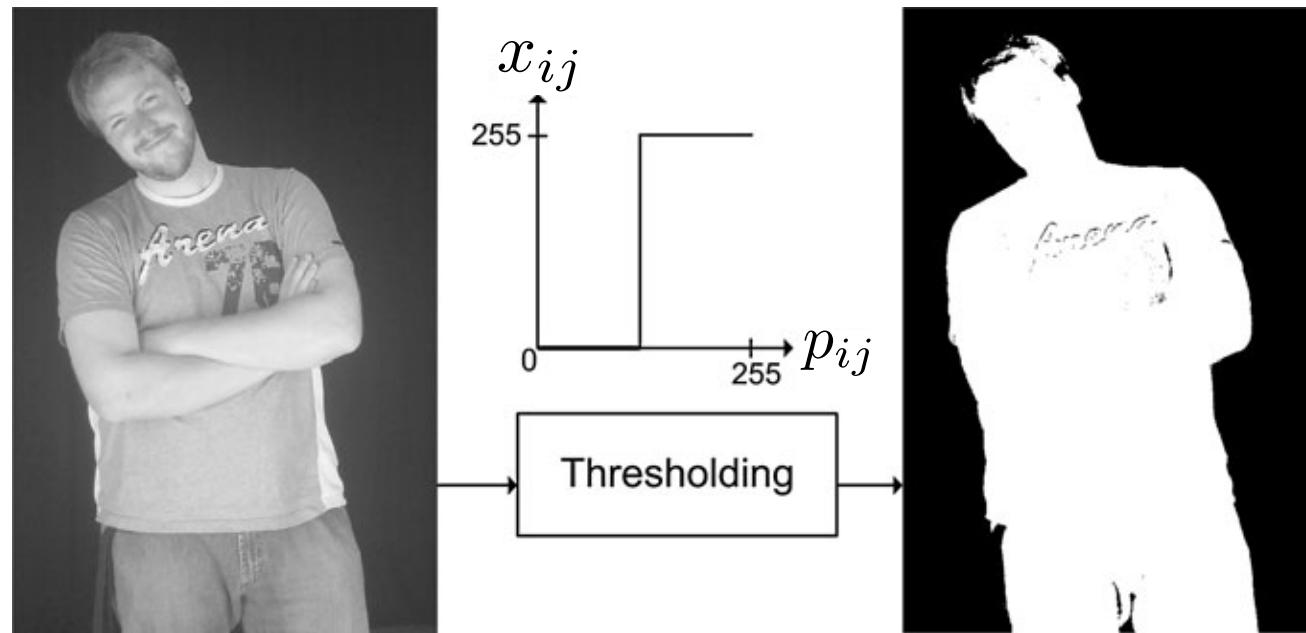
Equalization boosts the contrast more than whitened data – however, equalization is a more expensive operation.



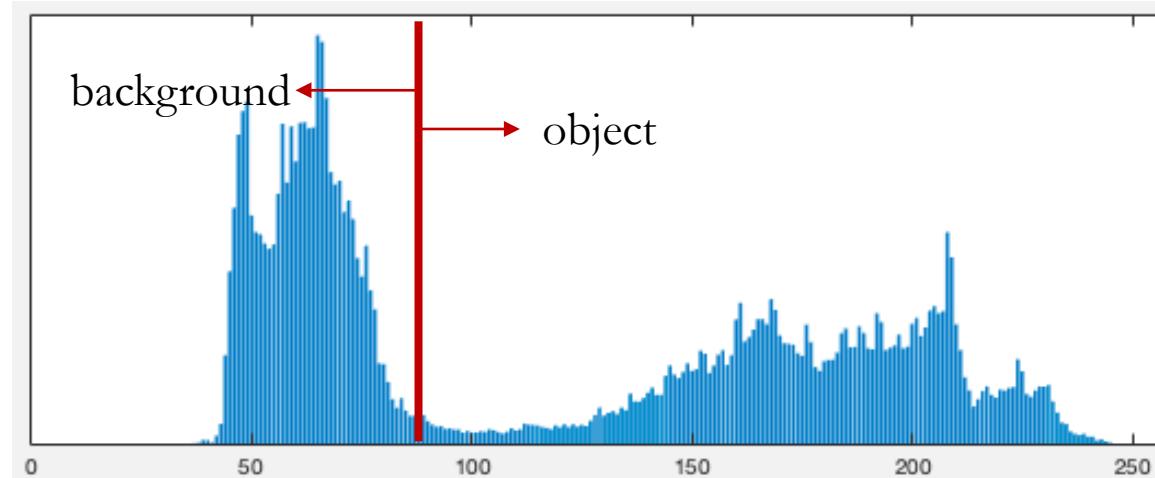
# Equalization vs. Stretching



# What are Histograms Good For?

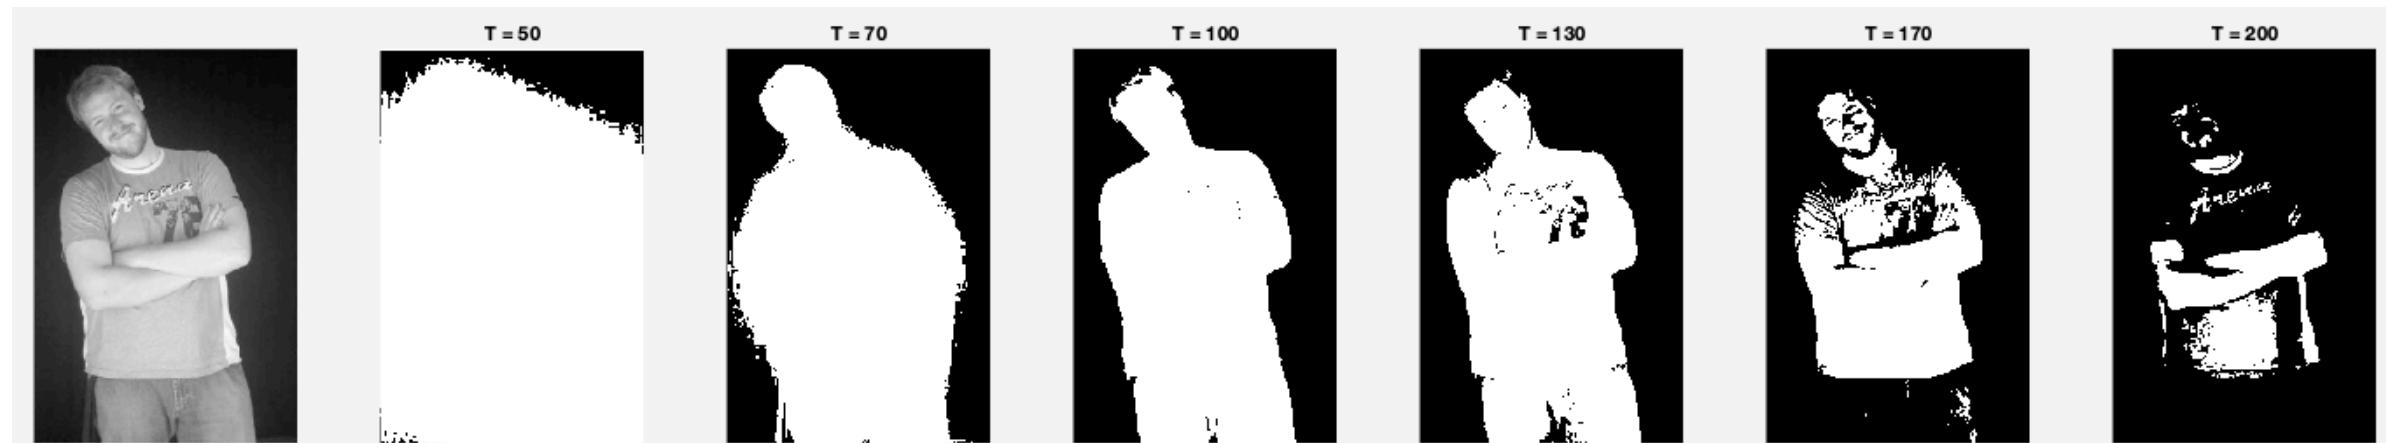


# Histogram Thresholding for FG/BG Separation

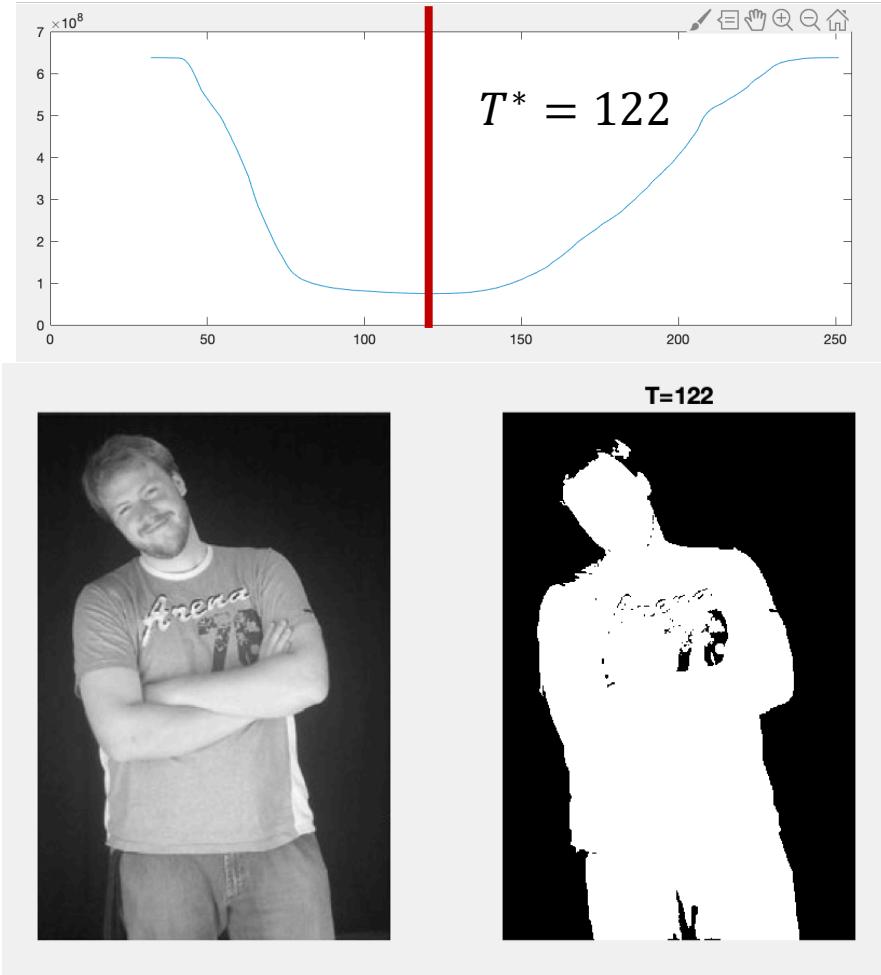


The threshold which separates the object from the background depends on the image. Can we choose the right threshold automatically?

If the image intensity histogram is bi-modal (two “humps”), then the threshold should sit somewhere in between the two modes.



# Otsu's Method – Automated Thresholding



$$T^* = \operatorname{argmin}_T w_1(T) \cdot \sigma_1^2(T) + w_2(T) \cdot \sigma_2^2(T)$$

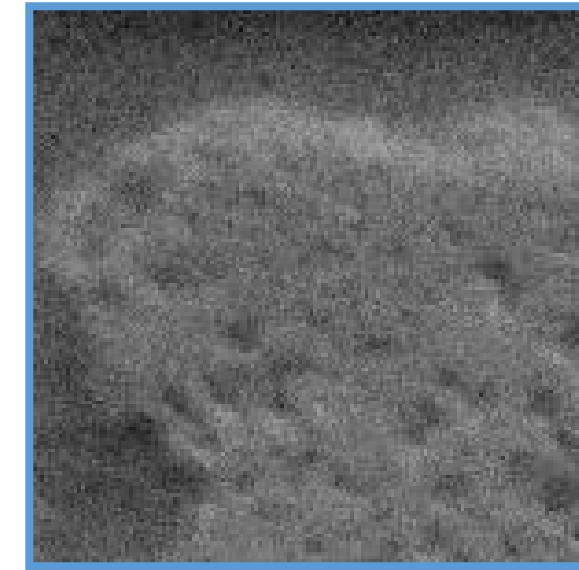
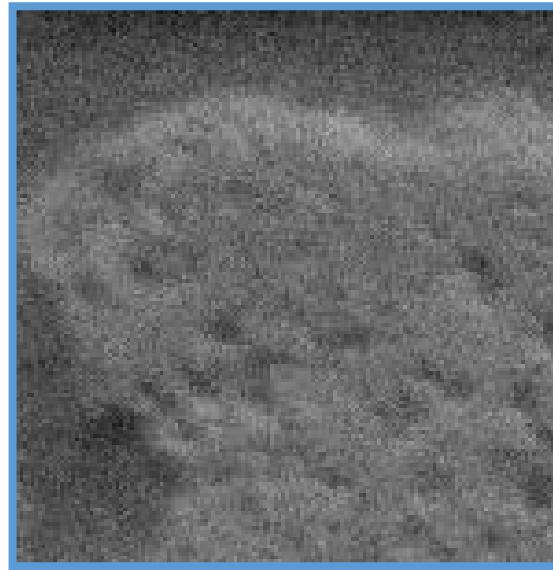
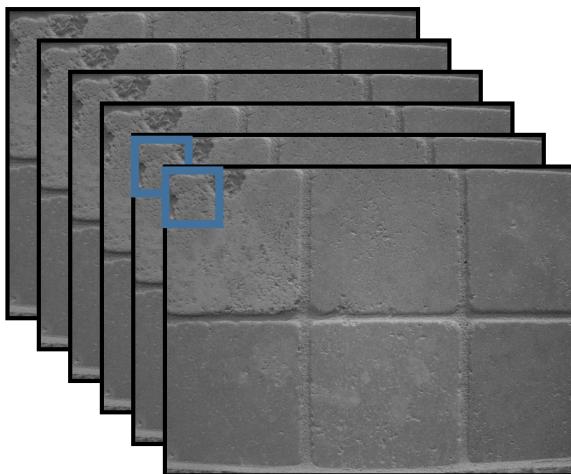
- optimal threshold  $T^*$  minimizes the sum of weighted variances of object and background.
- rationale:
  - correct threshold produces two narrow modes
  - incorrect threshold produces (at least one) wide mode
  - width of the mode is measured by the variance  $\sigma^2$
  - weighting by number of pixels in each mode
- $\sigma_1^2(T), \sigma_2^2(T)$  - variance of pixels less than or equal to and greater than threshold respectively
- $w_1(T), w_2(T)$  - number of pixels less than or equal to and greater than threshold

# Correlation Filters

Types of image noise; cross-correlation operation

Box & Gaussian smoothing; sharpening & template matching

# Motivation: Noise Reduction



Even multiple images of the **same static scene** will not be identical.  
Why not? Where does noise come from?

# Common types of noise

**Impulse & salt and pepper noise:** random occurrences of white (and black) pixels, usually due to defective sensor elements



Original



Impulse noise

**Gaussian noise:** variations in intensity drawn from a Gaussian distribution, due to inherent noise of the sensor

$$p_{ij} = \hat{p}_{ij} + \eta \quad \eta \sim \mathcal{N}(\mu_n, \sigma_n)$$

observation    ideal              noise  
        image

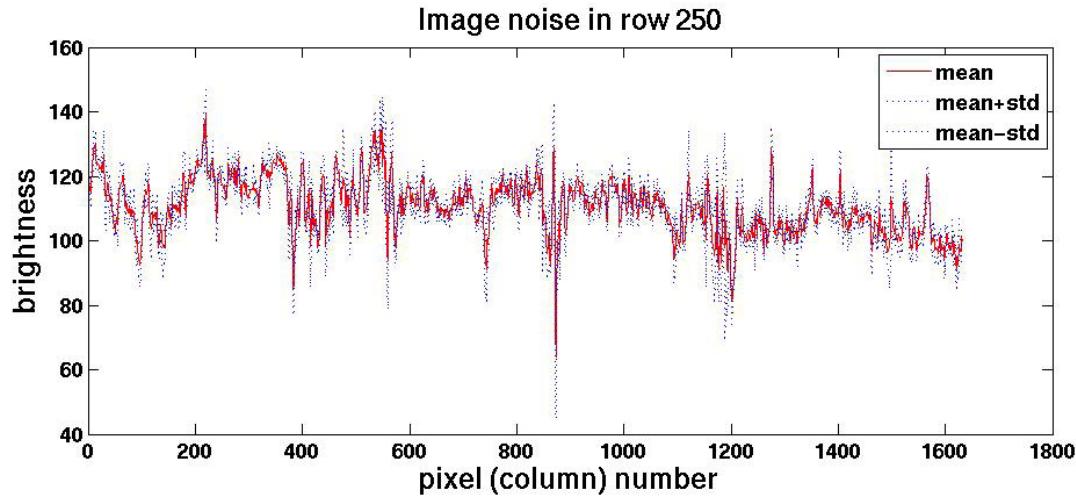
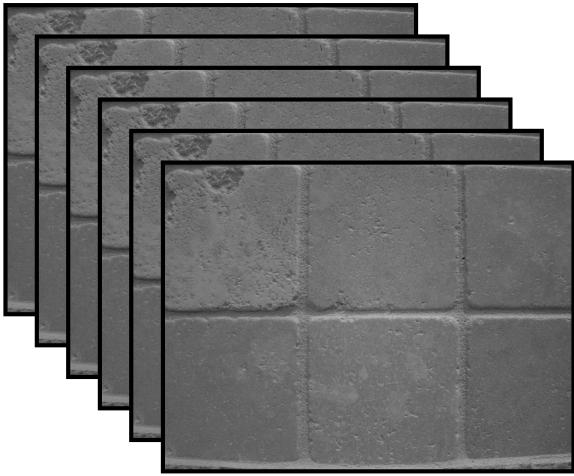


Salt and pepper noise



Gaussian noise

# Motivation: Noise Reduction



How could we reduce the noise, give an estimate of the true intensities?

What if there's only one image?

First attempt: Replace each pixel with average of all the values in its neighborhood.

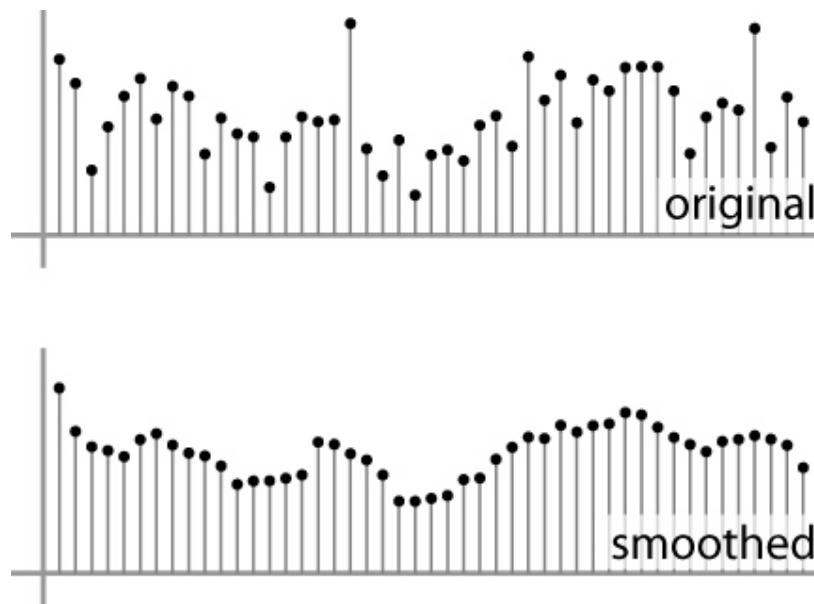
This assumes that

- (True) pixel values are similar to neighbours
- noise processes are independent from pixel to pixel.

# First Solution Attempt

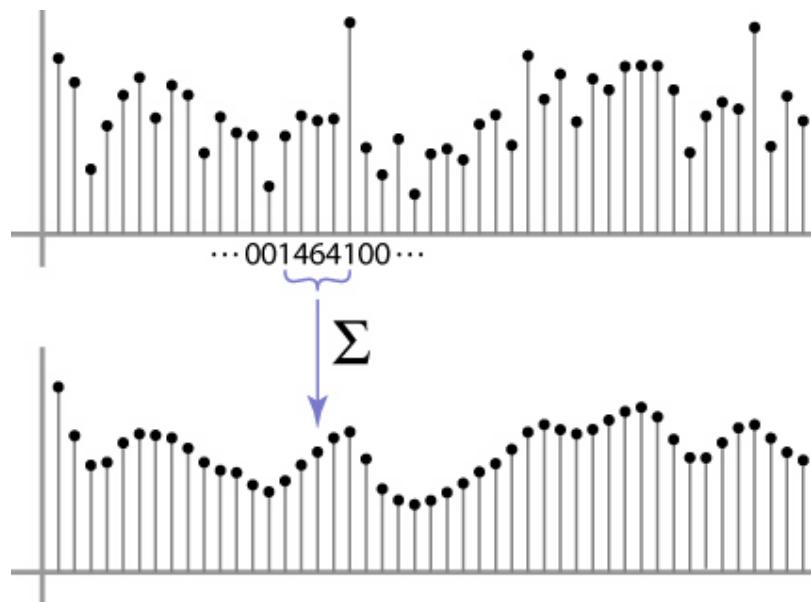
Replace each pixel with an average of all the values in its neighborhood

Moving average in 1D:



# Weighted Moving Average

- We can weight the moving average
  - Uniform weights:  $[1, 1, 1, 1, 1] / 5$
  - Non-uniform weights  $[1, 4, 6, 4, 1] / 16$



# Moving Average in 2D

P

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

X

	0	10	20	30	30	30	20	10	
0	20	40	60	60	60	40	20		
0	30	60	90	90	90	60	30		
0	30	50	80	80	90	60	30		
0	30	50	80	80	90	60	30		
0	20	30	50	50	60	40	20		
10	20	30	30	30	30	20	10		
10	10	10	0	0	0	0	0		

gives uniform weight  
to each pixel

$$x_{ij} = \frac{1}{2k+1}^2 \sum_{u=-k}^k \sum_{v=-k}^k p_{i+u,j+v}$$

sum all pixels in  
 $(2k+1) \times (2k+1)$   
neighborhood around  $p_{ij}$

# Correlation filtering

Generalize to allow different weights depending on neighboring pixel's relative position:

$$x_{ij} = \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u,j+v}$$

Can further generalize, no need to be square of  $(2k+1) \times (2k+1)$

Weight depends on neighbour's relative position wrt  $p_{ij}$

This is called cross-correlation, denoted  $\mathbf{X} = \mathbf{P} \otimes \mathbf{F}$

Cross-correlation is a linear operation.

Filter Kernel  
(stores the weights).

# Averaging filter (Box filter)

Which values belong in the kernel  $F$  for the moving average example?

$$\mathbf{P} \otimes \mathbf{F} = \mathbf{X}$$

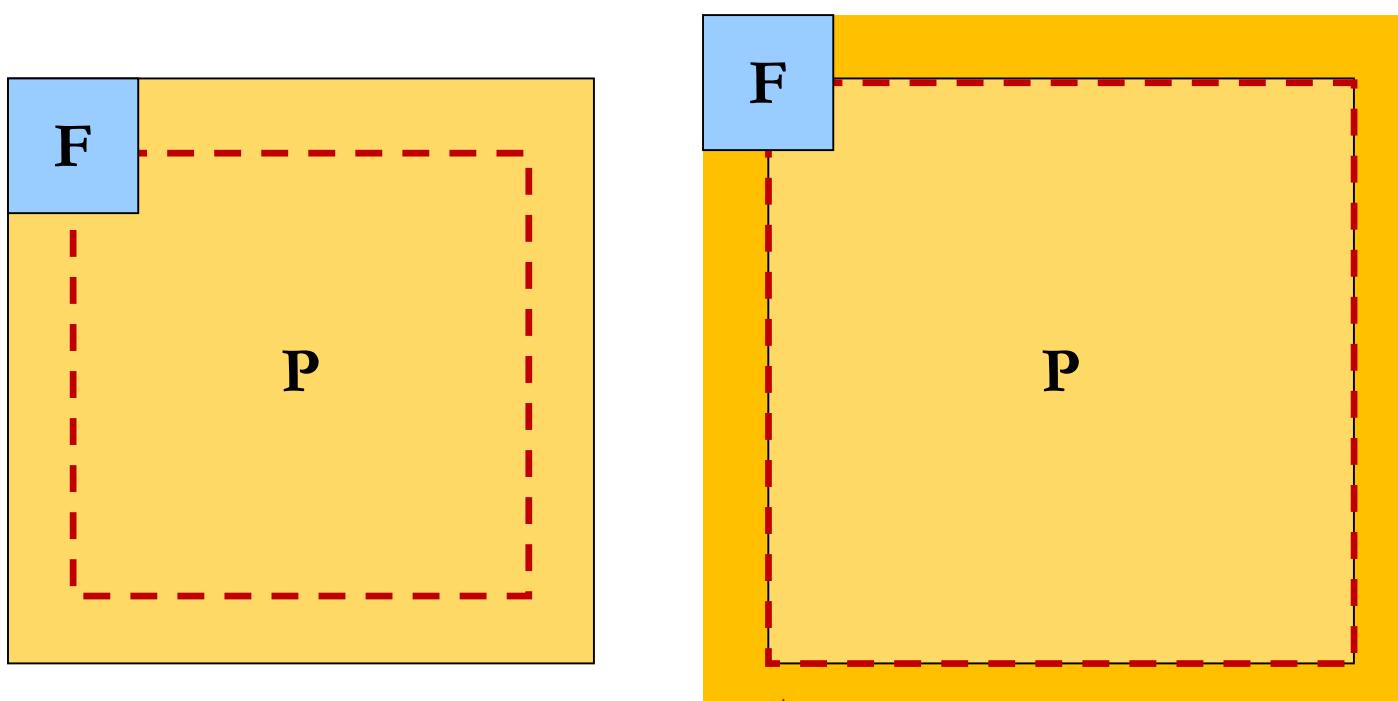
“box filter”

$\frac{1}{9}$

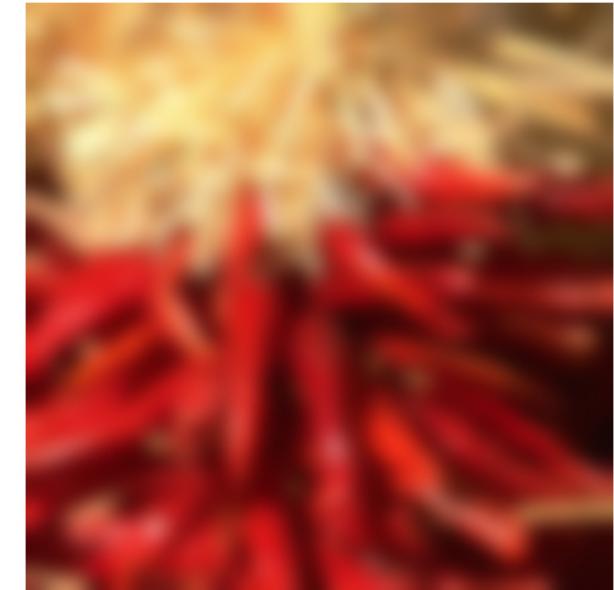
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

0	10	20	30	30	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

# Output Size & Boundaries



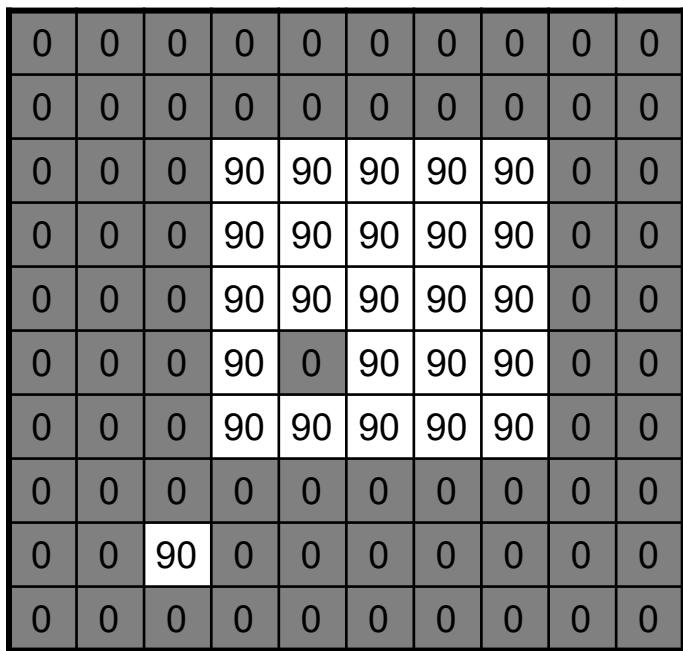
Filter window falls off  
the edge of the image  
- what goes in here?



- zero-padding
- wrap around
- copy edge
- reflect across edge

# Gaussian filter

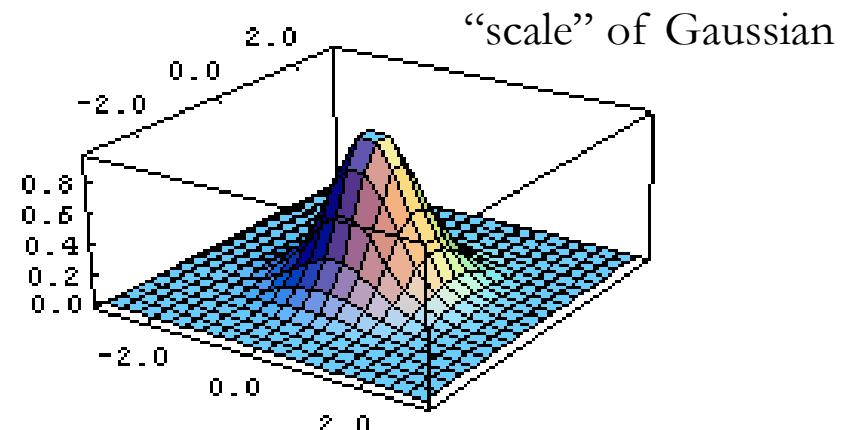
What if we want nearest neighboring pixels to have the most influence on the output?



$$\frac{1}{16}$$

1	2	1
2	4	2
1	2	1

$$f_{uv} = \frac{1}{2\pi\underline{\sigma}^2} e^{-\frac{u^2+v^2}{\underline{\sigma}^2}}$$

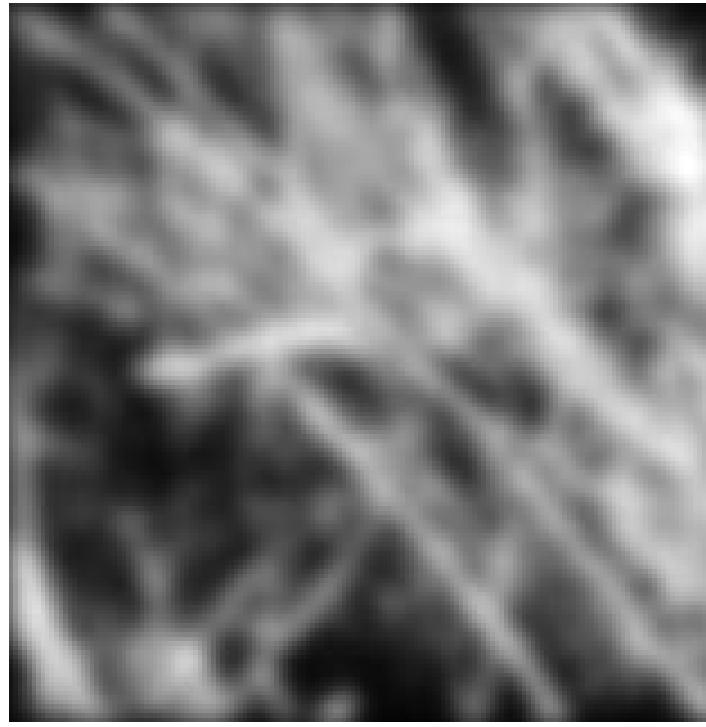


Both the box filter and the Gaussian filter has a smoothing effect.

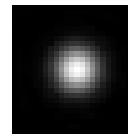
# Smoothing with Gaussian vs. Box



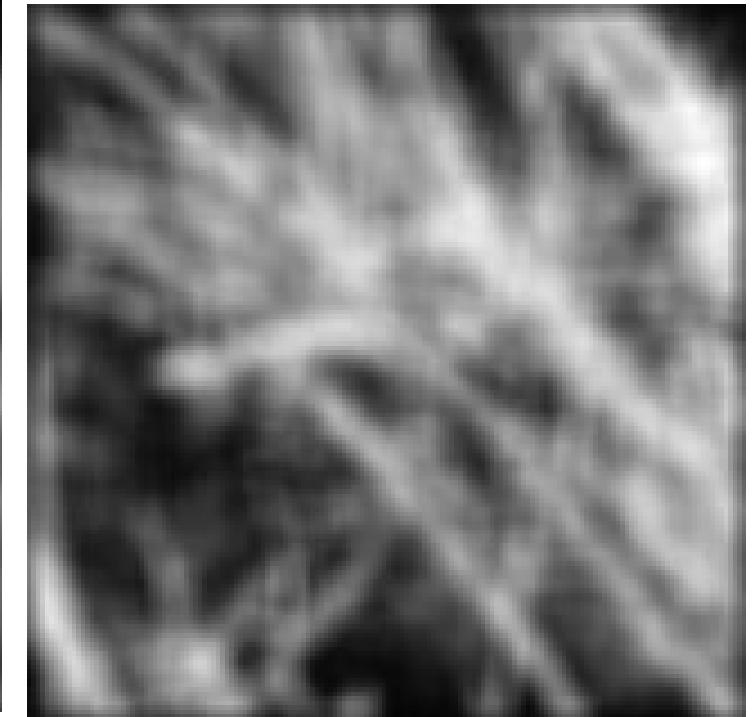
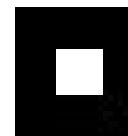
original



filtered  
(Gaussian kernel)



depicts filter kernels:  
white = high value  
black = low value



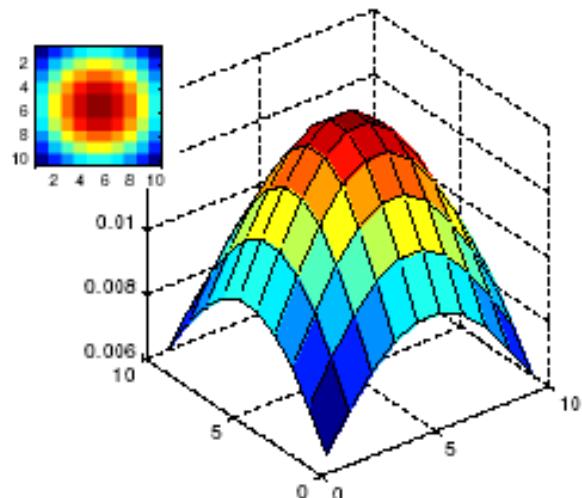
filtered  
(Box kernel)

# Gaussian filters

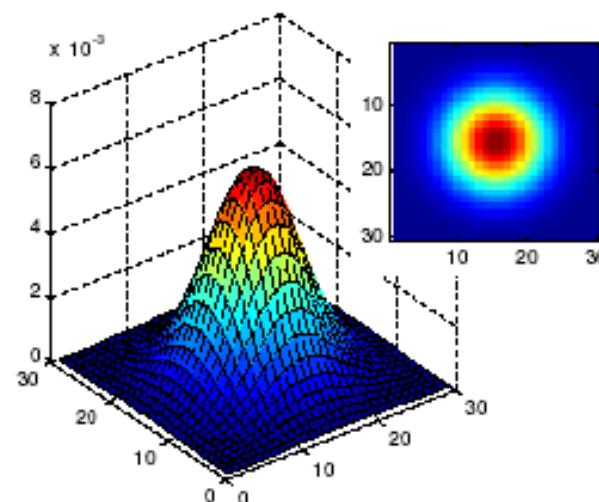
Which parameters matter?

**Size of kernel?**

Note, Gaussian function has infinite support, but discrete filters use finite kernels



$\sigma = 5$  with  $10 \times 10$  kernel

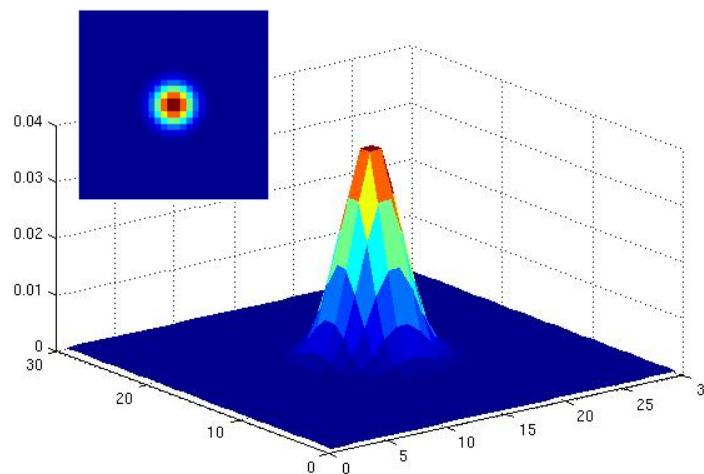


$\sigma = 5$  with  $30 \times 30$  kernel

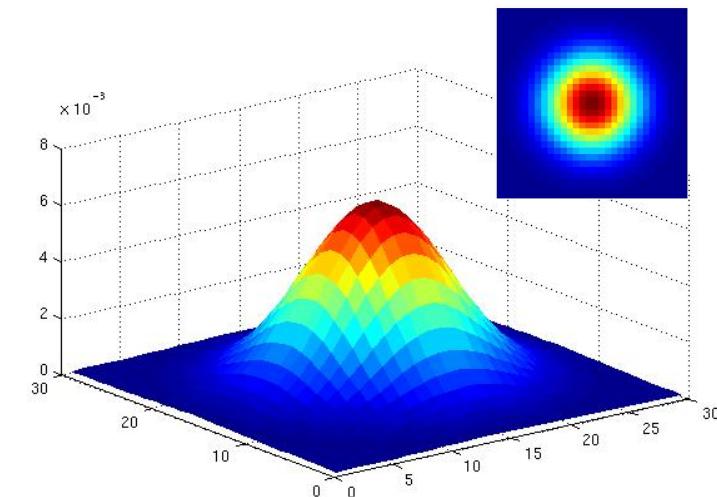
# Gaussian filters

Which parameters matter?

**Variance of Gaussian ( $\sigma$ ):** determines extent of smoothing



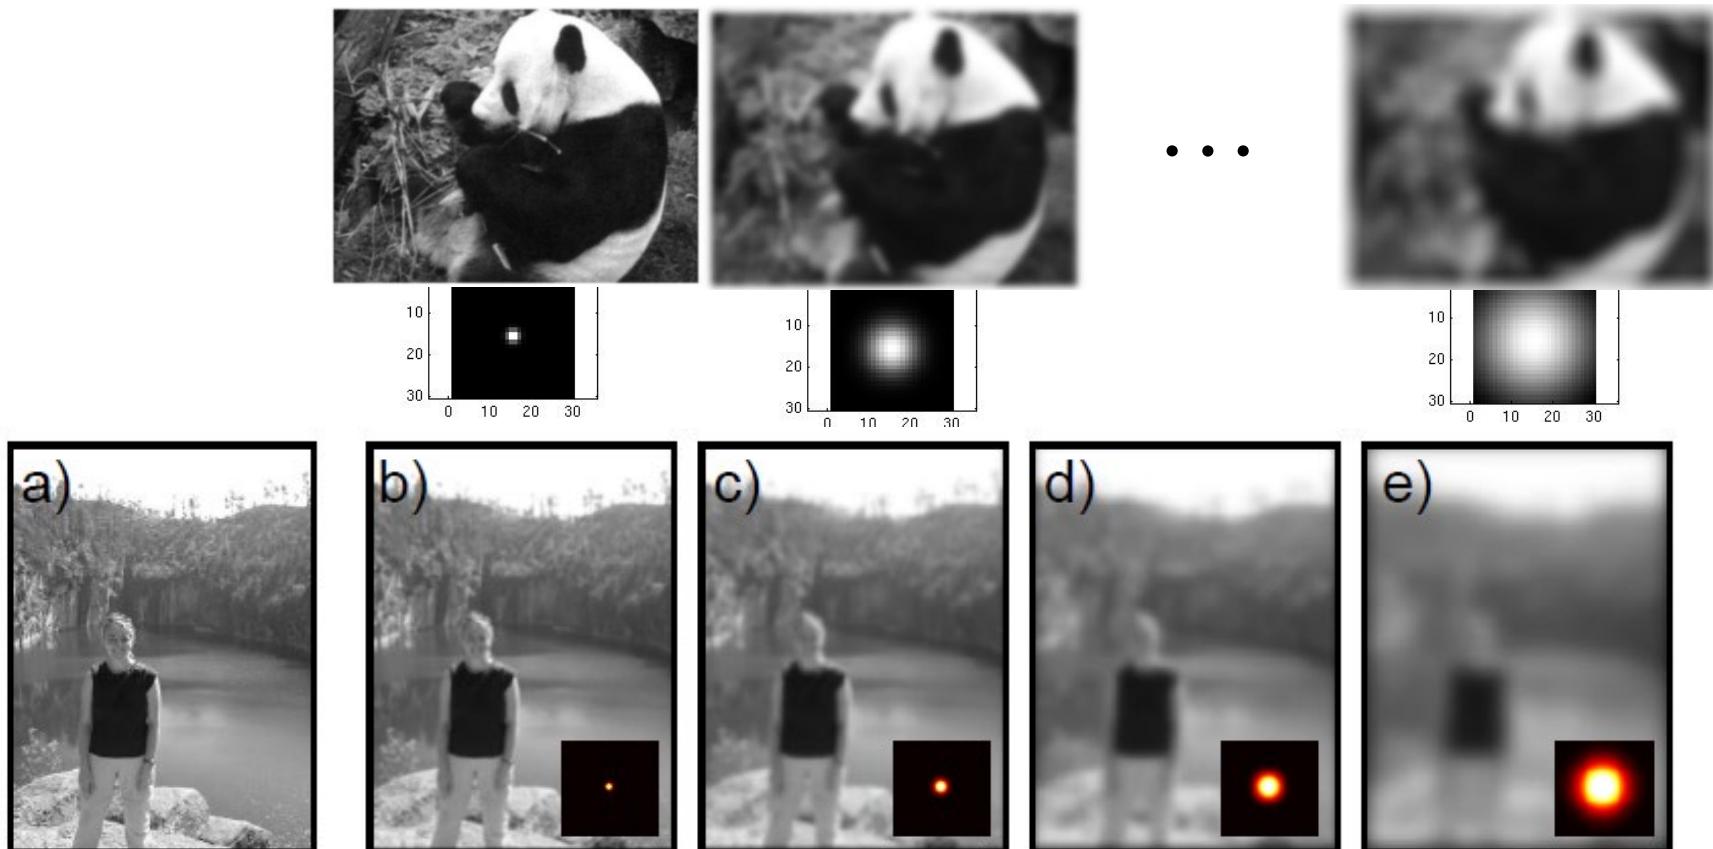
$\sigma = 2$  with  $30 \times 30$  kernel



$\sigma = 5$  with  $30 \times 30$  kernel

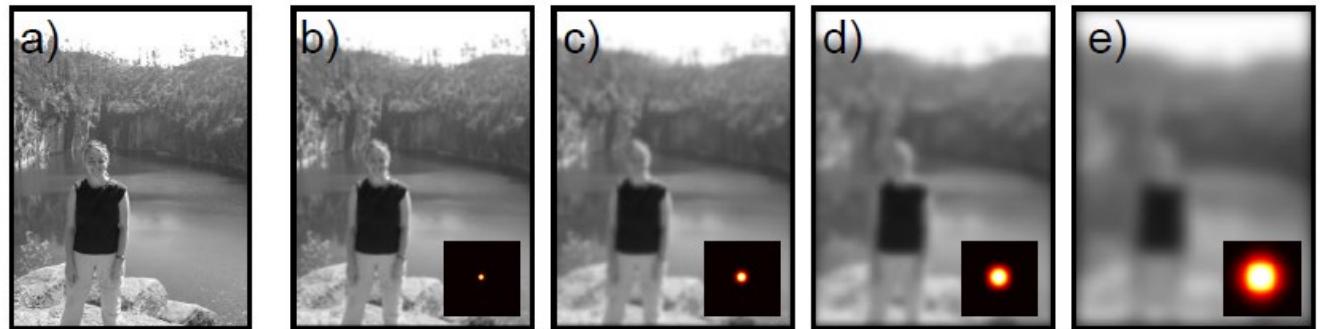
# Smoothing with a Gaussian

$\sigma$  is the “scale” of the Gaussian kernel, and controls the amount of smoothing.



# Gaussian Noise vs. Gaussian Smoothing

- Gaussians are used as a filter kernel applied for blurring or smoothing
- Gaussians are also used to model the noise processes.



$$p_{ij} = \hat{p}_{ij} + \eta \quad \eta \sim \mathcal{N}(\mu_n, \sigma_n)$$

observation  
ideal  
image

noise

In real life, the “ideal image” is not obtainable. We show (manually) corrupted versions to illustrate the impact of additive Gaussian noise.

- You can use Gaussian smoothing (or any other spatial smoothing kernel) to reduce the effects of Gaussian noise ☺

# Practice with Correlation Filters



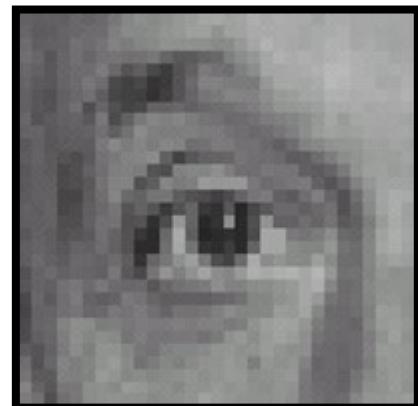
Original



0	0	0
0	1	0
0	0	0



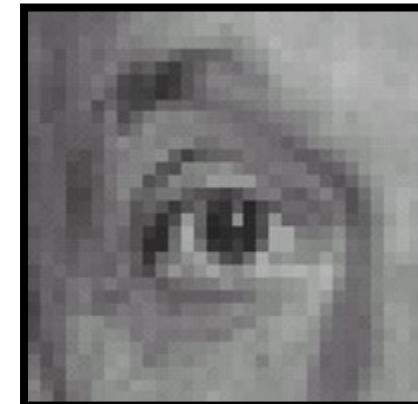
Filtered (no change)



Original

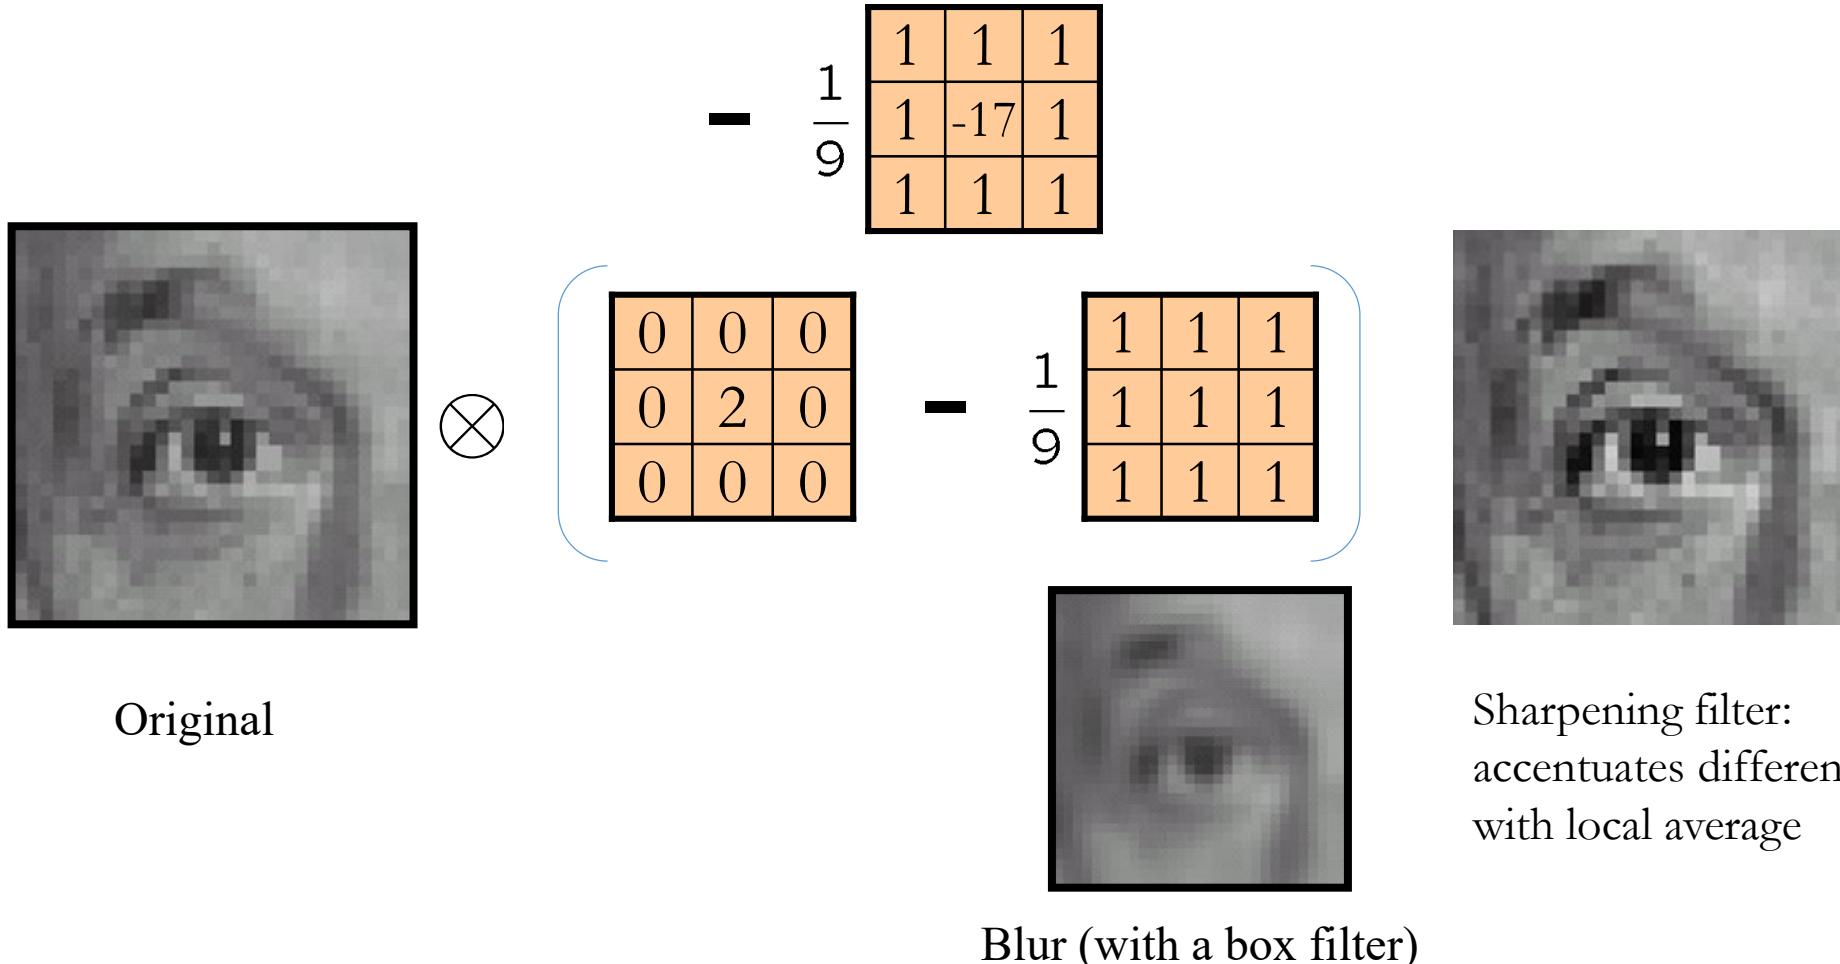


0	0	0
0	0	1
0	0	0

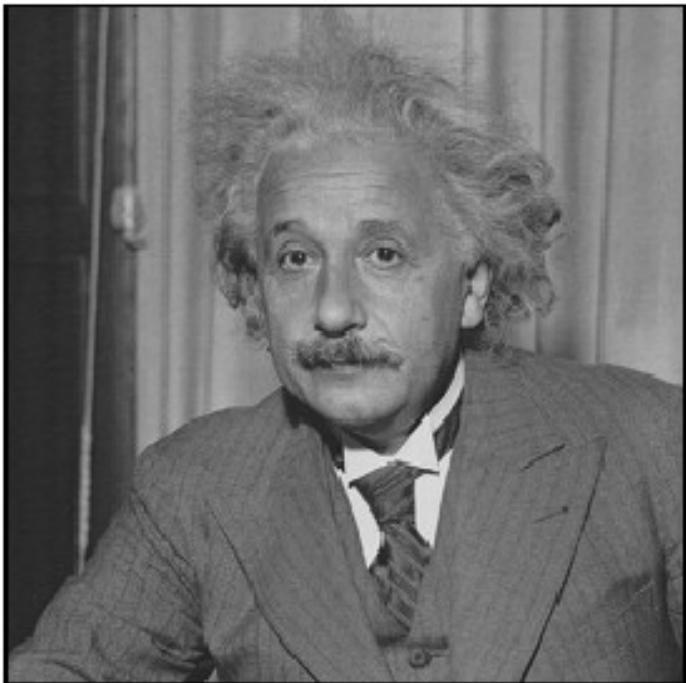


Filtered (shifts to the left)

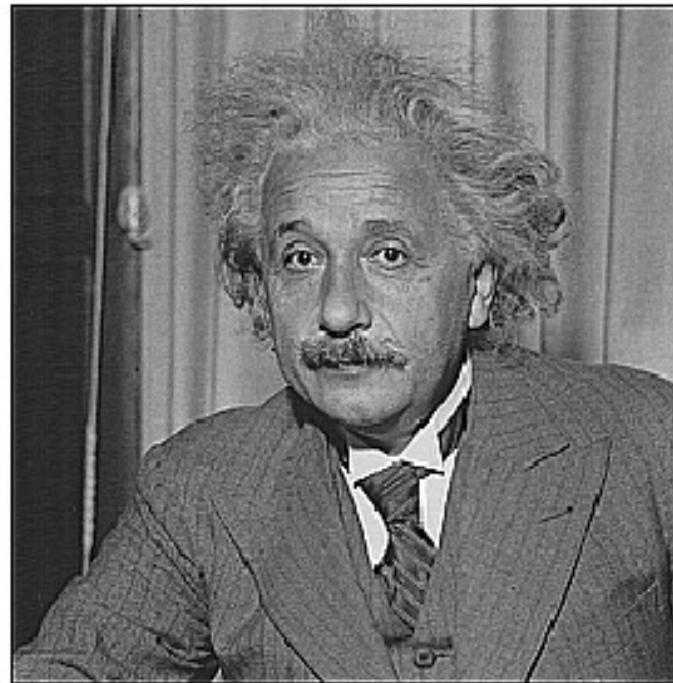
# Practice with Correlation Filters



# Understanding Sharpening



**before**



**after**

- do nothing for flat areas
- stresses intensity peaks and differences with respect to the surroundings

# Understanding Sharpening

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

Response if region is “flat”?

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

No change

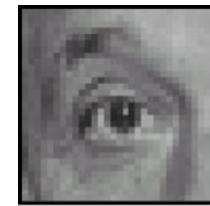
0	0	1	1	1
0	0	1	1	1
0	0	1	1	1
0	0	1	1	1
0	0	1	1	1

Region has an edge?

0	-2	14	12	14
0	-3	12	9	12
0	-3	12	9	12
0	-3	12	9	12
0	-2	14	12	14

Edge gets emphasized  
02. Point Processing & Filtering

input



filter

0	0	0
0	2	0
0	0	0

$$- \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

output



sharpening

0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0
0	0	0	0	0

Region has noise?

0	0	0	0	0
0	-1	-1	-1	0
0	-1	17	-1	0
0	-1	-1	-1	0
0	0	0	0	0

Noise also emphasized ☹

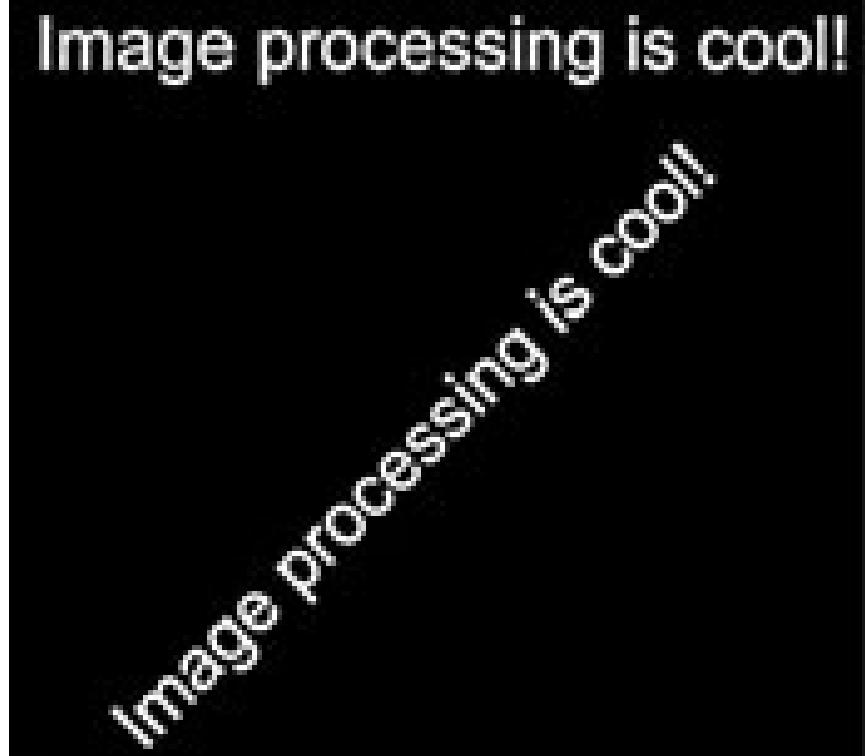
# Template Matching

processing

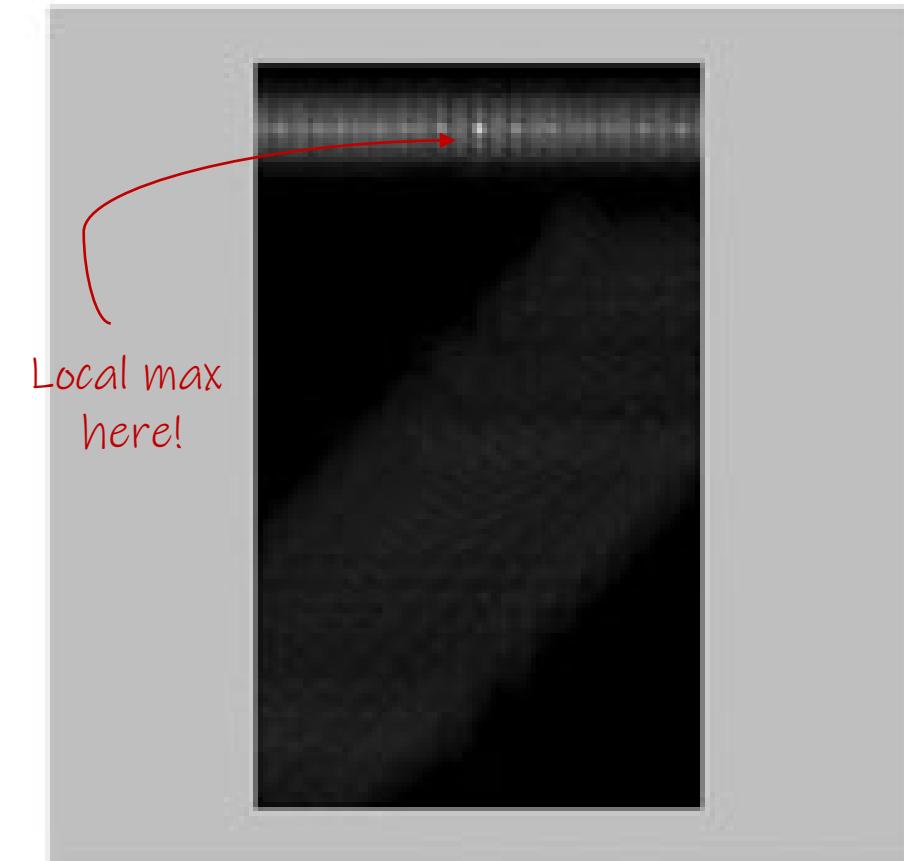
"filter kernel" i.e.  
matrix of numbers

Input image

Matrix of numbers



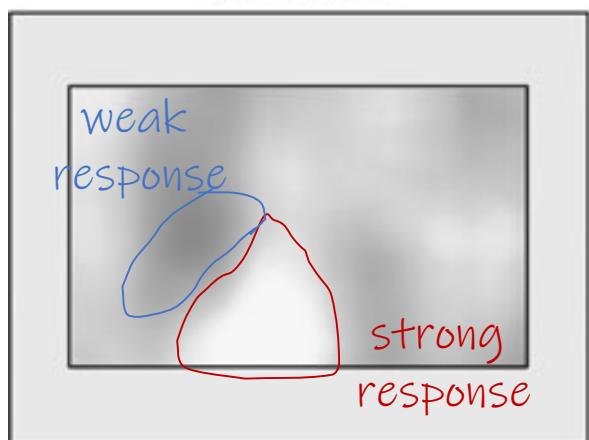
Correlation result



# Normalized Cross-Correlation



Normalized cross correlation



$$x_{ij} = \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u, j+v}$$

Scale to sum  
up to 1

[0-255]!

$$x_{ij} = \frac{1}{|\mathbf{F}| |\mathbf{w}_{ij}|} \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u, j+v}$$

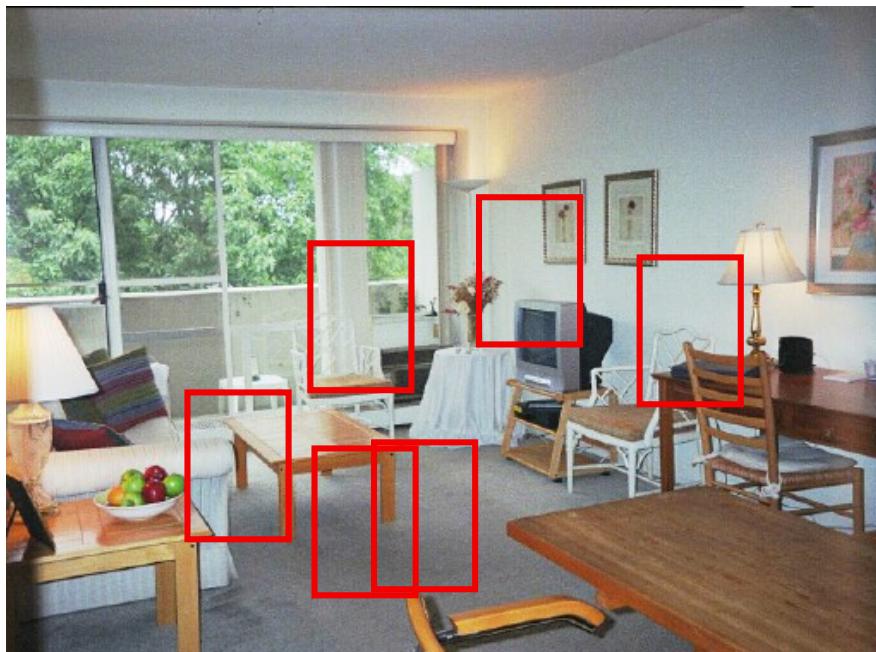
Normalized version scales with respect to both the filter and the corresponding input window

$|\bullet|$  indicates magnitude of kernel, i.e. square root of sum of squares of all elements in kernel.

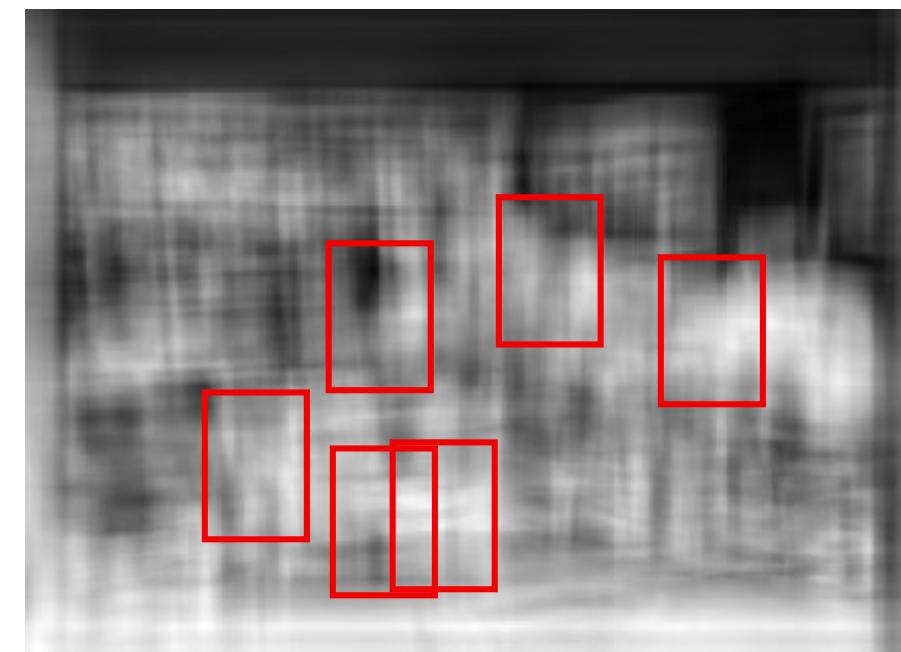
Local max here!

# Template Matching for Detecting Objects?

Find the chair in this image



Simple template matching doesn't cut it...



A “popular method is that of template matching, by point to point correlation of a model pattern with the image pattern. These techniques are inadequate for three-dimensional scene analysis for many reasons, such as occlusion, changes in viewing angle, and articulation of parts.” Nevatia & Binford, 1977.

# Convolution Filtering

Convolution operation;

Convolution vs. cross-correlation

# Filtering an impulse signal

What is the result of filtering an impulse signal (image) with an arbitrary kernel?

$$\mathbf{P} \otimes \mathbf{F} = \mathbf{X}$$

The diagram illustrates the convolution process between an input image  $\mathbf{P}$  and a filter  $\mathbf{F}$  to produce the output image  $\mathbf{X}$ .

**Input Image ( $\mathbf{P}$ ):**

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

**Filter ( $\mathbf{F}$ ):**

a	b	c
d	e	f
g	h	i

**Output Image ( $\mathbf{X}$ ):**

		<b>i</b>	<b>h</b>	<b>g</b>			
		<b>f</b>	<b>e</b>	<b>d</b>			
		<b>c</b>	<b>b</b>	<b>a</b>			

The diagram shows the receptive field of the central pixel 'i' in the output image, which covers the entire input image  $\mathbf{P}$ . The filter  $\mathbf{F}$  is applied to the central pixel '1' of the input image  $\mathbf{P}$ , resulting in the output value 'i'.

# Convolution

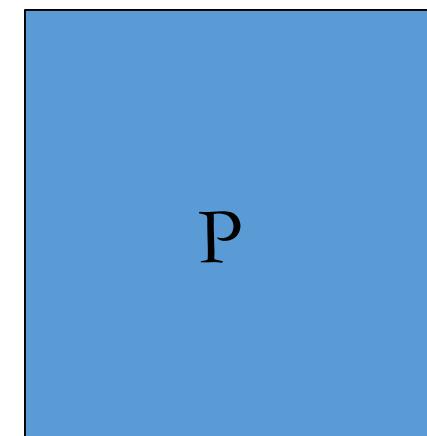
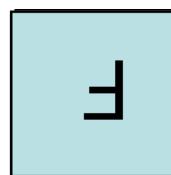
Flip the kernel in both dimensions, then apply cross-correlation

$$x_{ij} = \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i-u, j-v}$$

Subtraction of  $u, v$  indices!

$$\mathbf{X} = \mathbf{F} * \mathbf{P}$$

convolution operator



# Convolution vs. Cross-correlation

Convolution 
$$x_{ij} = \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i-u, j-v} \quad \mathbf{X} = \mathbf{P} * \mathbf{F}$$

Cross-correlation 
$$x_{ij} = \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u, j+v} \quad \mathbf{X} = \mathbf{P} \otimes \mathbf{F}$$

# Properties of Convolution

Commutative:

$$f * g = g * f$$

Associative:

$$(f * g) * h = f * (g * h)$$

Distributes over addition:

$$f * (g + h) = (f * g) + (f * h)$$

Scalars factor out:

$$\lambda f * g = f * \lambda g = \lambda(f * g)$$

Identity:

$$f * e = f, \text{ where } e = [\dots, 0, 0, 1, 0, 0, \dots].$$

Shift Invariant:

behaves the same everywhere, i.e. output value depends only on the pattern and not position of neighbourhood

# Convolution vs. Correlation

- convolution is the “standard” we refer to when we talk about filtering unless otherwise explicitly stated
- for images, convolution and correlation are often used interchangeably since the two operations are equivalent if the filter kernels are symmetric
- symmetric: if “flipping” of kernel does not have any effect, e.g.
  - Gaussian
  - Box kernel

# Non-Linear Filters

Median filtering, denoising impulse noise

# Non-Linear Filters

- Correlation and convolution are linear operations
- Non-linear filters perform non-linear operations, e.g.
  - Median
    - Median filtering is commonly used for removing impulse and salt & pepper noise
  - Min/Max
    - The max operation is used in max-pooling for deep neural networks.
  - Morphology
- We will take a closer look at the median filter

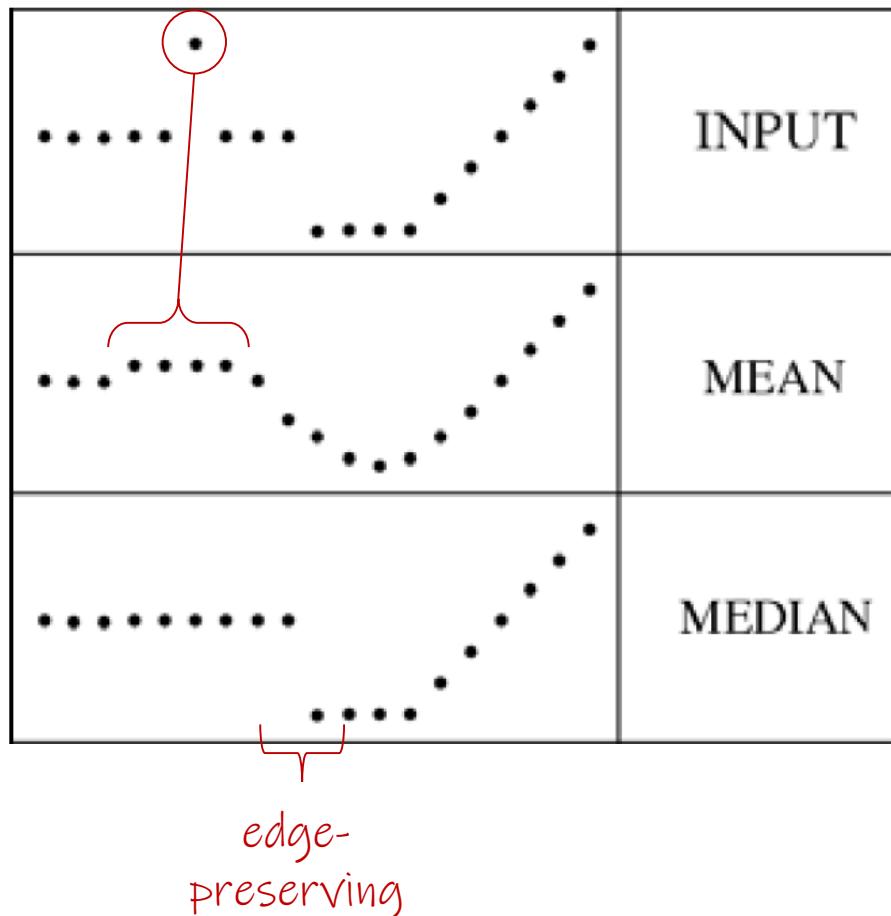
# Denoising: Median Filter



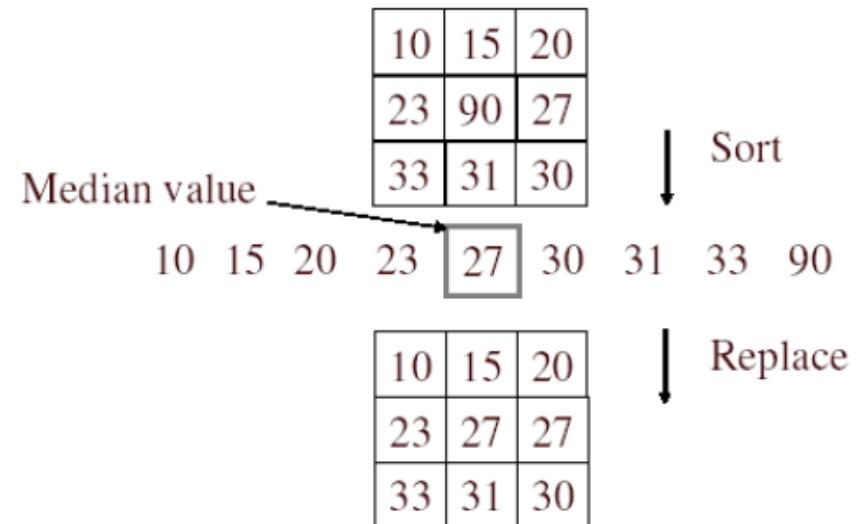
Salt and pepper noise



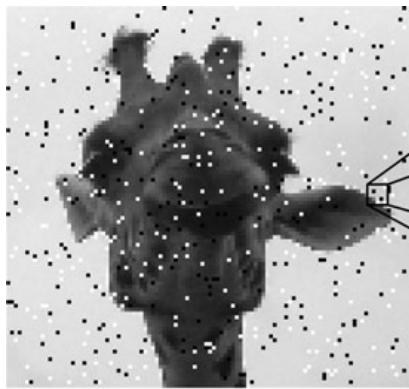
Impulse noise



- No new pixel values introduced
- Removes spikes: good for impulse, salt & pepper noise



# Median Filter Outputs

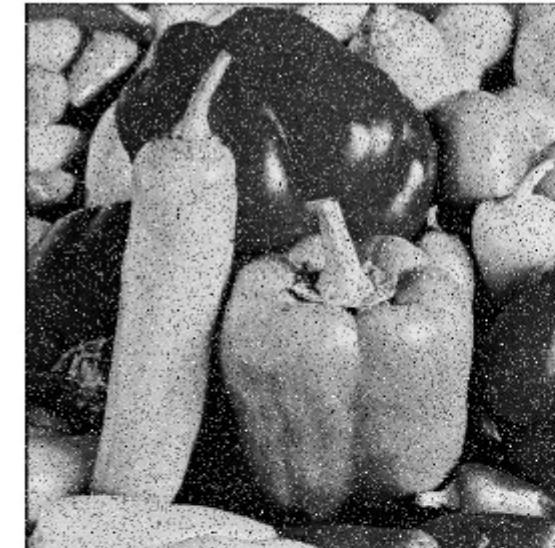


Mean filtered

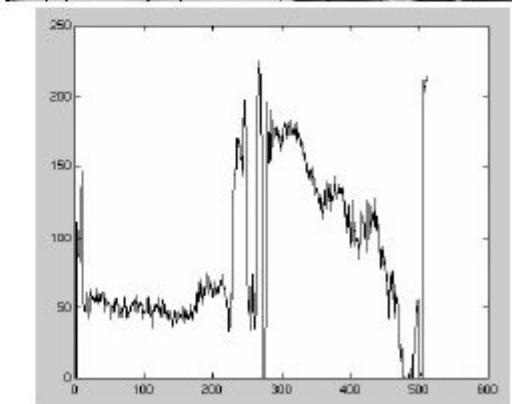
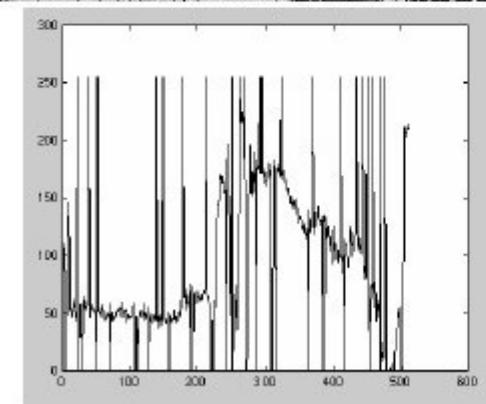
Median filtered

205	204	204	206	255
206	0	208	206	206
201	199	205	206	209
61	128	213	0	205
59	65	255	206	255

Salt and pepper noise



Median filtered



Plots of a scanline.

# Summary

## 1. Point-wise Processing

- Brightness, contrast, normalization, non-linear mappings
- Histograms: stretching, equalization, segmentation

## 2. Linear Filtering

- Achieves many effects: denoising, sharpening, moving, template matching
- (Cross-)Correlation:
  - weighted summation of input window, weights stored in filter / kernel
  - normalized version for template matching
- Convolution
  - Flips kernel before applying cross-correlation

## 3. Non-linear filters

- Median filter for smoothing salt & pepper, impulse noise