

Gradients & Edges

CS 4243 Computer Vision & Pattern Recognition

Angela Yao

Recap & Outline

Last Week

- Linear filtering, cross-correlation vs. convolution operations
- Denoising, sharpening, template matching
- Gaussian & Laplacian image pyramids
- Non-linear filtering, e.g. median filters

Today's Lecture

- Motivation for studying image edges
- Derivative filters to extract gradients
- Need for smoothing
- Canny edge detector

Why Study Edges?




Figure from J. Shotton et al., PAMI 2007

- Human visual system is sensitive to edges
- Edges compactly convey salient information


Image Edges Convey Information

reflectance changes,
textures & appearances



depth discontinuities &
object boundaries

change in surface
orientations & shape


shadows



Why Extract Edges?



Credit: Attneave



Credit: Jitendra Malik

Resilient to lighting and color → useful for recognition

Cue for shape and geometry
→ useful for recognition,
3D understanding

Image Gradients

Definition, magnitude & orientation, Sobel filtering

What are Image Edges or Gradients?






Image edges are sharp discontinuities in the intensity.

Derivatives are large at discontinuities. Edges correspond to extrema of derivative.


image



intensity along horizontal scanline






1st derivative



Understanding Image Gradients

$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$ The gradient is a vector that points in the direction of most rapid change in intensity.



The **gradient direction** (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

The **edge strength** is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

Images are discrete. How to take (partial) derivatives of discrete signals?

Take finite differences.

Finite differences

Recall: definition of a derivative (with a forward difference)

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

Alternative with central difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+0.5h) - f(x-0.5h)}{h}$$

Discrete version: remove limit, set $h = 2$

$$f'(x) = \frac{f(x+1) - f(x-1)}{2}$$

What filter kernel does this correspond to?

1D derivative filter

correlation	<table><tr><td>-1</td><td>0</td><td>1</td></tr></table>	-1	0	1
-1	0	1		
convolution	<table><tr><td>1</td><td>0</td><td>-1</td></tr></table>	1	0	-1
1	0	-1		

The Sobel filter

Horizontal Sobel filter:

Horizontal gradients \rightarrow vertical lines

1	0	-1
2	0	-2
1	0	-1

=


1
2
1

*

1	0	-1
---	---	----

1D derivative
filter

What filter is this? *blurring*



Vertical Sobel filter:

Vertical gradients \rightarrow horizontal lines


1	2	1
0	0	0
-1	-2	-1

=


1	2	1
---	---	---

*

1
0
-1




Sobel filter example



original



horizontal Sobel filter



vertical Sobel filter

Comparison of Gradient Operators



Sobel



1	0	-1
2	0	-2
1	0	-1
1	2	1
0	0	0
-1	-2	-1

Scharr



3	0	-3
10	0	-10
3	0	-3
3	10	3
0	0	0
-3	-10	-3

Prewitt



1	0	-1
1	0	-1
1	0	-1
1	1	1
0	0	0
-1	-1	-1

Roberts



0	1
-1	0
1	0
0	-1

Most well-known, good tradeoff between simplicity and efficiency.

Computing Image Gradients

1. Select your favorite gradient filters.

$$\mathbf{S}_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$\mathbf{S}_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

2. Convolve with the image \mathbf{f} to compute image gradient components.

$$\frac{\partial \mathbf{f}}{\partial x} = \mathbf{S}_x * \mathbf{f} \qquad \frac{\partial \mathbf{f}}{\partial y} = \mathbf{S}_y * \mathbf{f}$$

3. Compute gradient orientation and magnitude.

$$\nabla \mathbf{f} = \left[\frac{\partial \mathbf{f}}{\partial x}, \frac{\partial \mathbf{f}}{\partial y} \right]$$




gradient

$$\theta = \tan^{-1} \left(\frac{\partial \mathbf{f}}{\partial y} / \frac{\partial \mathbf{f}}{\partial x} \right)$$

orientation

$$\|\nabla \mathbf{f}\| = \sqrt{\left(\frac{\partial \mathbf{f}}{\partial x} \right)^2 + \left(\frac{\partial \mathbf{f}}{\partial y} \right)^2}$$

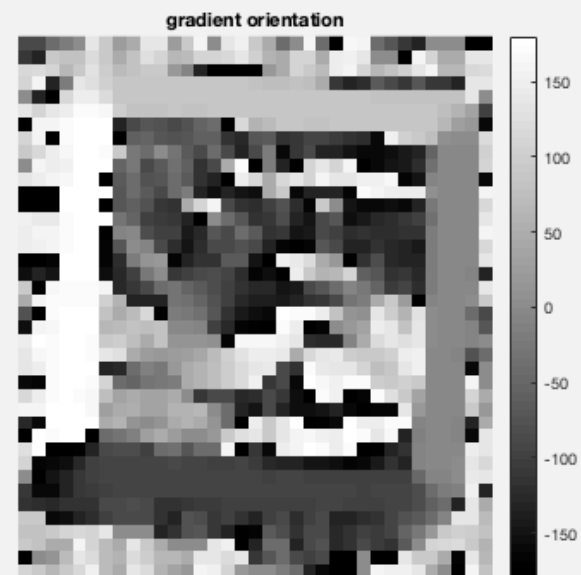

magnitude



gradient $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

magnitude $\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$


orientation $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$




DoG & Laplace Filters

Derivative of Gaussians, Laplacian of Gaussians

How do you find the edge in this signal?




intensity plot



Using a derivative filter:

Where is the
extrema here?



Noise is high-frequency.
Differentiation accentuates noise


$$\frac{d \sin \omega x}{dx} = \omega \cos \omega x$$

Multiplicative factor
proportional to frequency!


Differentiation is very sensitive to noise

When using derivative filters, it is critical to blur first!


f input




h Gaussian



$h * f$ blurred



$\frac{\partial}{\partial x}(h * f)$ derivative of blurred



Where is the edge?


Look for peaks in

$$\frac{\partial}{\partial x}(h * f)$$

Trade Off Noise vs. Edge Localization



Image with Edge



Edge Location






Image + Noise



W/out smoothing, edge and noise is picked up by derivative




Smoothing removes the noise, but also blurs the edge

Derivative of Gaussian Filter

Differentiation Property of Convolution: $\frac{\partial}{\partial x}(h * f) = (\frac{\partial}{\partial x}h) * f$


f

input




$\frac{\partial}{\partial x}h$

derivative of
Gaussian



$(\frac{\partial}{\partial x}h) * f$

output (same as
before)




Derivative of Gaussian Filter

Differentiation is a convolution operation.

$$\frac{\partial}{\partial x}(h * f) = \left(\frac{\partial}{\partial x}h\right) * f = (h * g) * f$$


Which kernel?






*

$$\begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

=





What should we set as σ for the Gaussian?




Smoothing with a Gaussian

σ is the “scale” of the Gaussian kernel and determines the extent of smoothing.



...



Effect of σ on Derivatives

Depending on the applied σ , different structures appear. Larger σ detects larger-scale edges, smaller σ detects finer edges.









FIGURE 5.3: The scale (i.e., σ) of the Gaussian used in a derivative of Gaussian filter has significant effects on the results. The three images show estimates of the derivative in the x direction of an image of the head of a zebra obtained using a derivative of Gaussian filter with σ one pixel, three pixels, and seven pixels (**left to right**). Note how images at a finer scale show some hair, the animal's whiskers disappear at a medium scale, and the fine stripes at the top of the muzzle disappear at the coarser scale.


Which σ Should We Use?

— true signal
— noisy observation

$\frac{\partial}{\partial x}(h * f)$ derivative of
smoothed signal



Large $\sigma \rightarrow$ Good detection (high SNR)
Poor localization



Small $\sigma \rightarrow$ Poor detection (low SNR)
Good localization

Gets difficult to
find the correct
maximum
corresponding to
the edge

Laplace filter

A second derivative filter that can also be approximated with finite differences.

first-order
finite difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$$



1D derivative filter

1	0	-1
---	---	----

second-order
finite difference

$$f''(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - 2f(x) + f(x - h)}{h^2}$$



1D Laplace filter

1	-2	1
---	----	---

Discrete
Laplacian
of x,y

$$\begin{aligned} \nabla^2 f(x, y) = & f(x + 1, y) + f(x - 1, y) \\ & + f(x, y + 1) + f(x, y - 1) - 4f(x, y) \end{aligned}$$




2D Laplace filter

0	1	0
1	-4	1
0	1	0

Sharpening with Laplacians



A: original



B: filtered with

0	1	0
1	-4	1
0	1	0

Derivative filter: highlights transitions. Mostly black because negative values are clipped to 0 by display.



A-B: sharpening effect

Add original image to sharpened details. Use a subtraction sign because centre coefficient of kernel (-4) is negative!



A-filtering with


1	1	1
1	-8	1
1	1	1

Extension of Laplacian kernel with diagonal terms.


Laplacian of Gaussian (LoG) filter

As with derivative, we can combine Laplace filtering with Gaussian filtering


input



Laplacian of
Gaussian




output




“zero crossings” at edges

2D Gaussian filters




$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

Gaussian



$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$

Derivative of Gaussian



Laplacian
Operator

$$\nabla^2 h_{\sigma}(u, v)$$

Laplacian of Gaussian

Laplace and LoG filtering examples

Also sensitive to noise



Laplace filtering

Smoothing w/ Gaussian helps




Laplacian of Gaussian filtering




Derivative of Gaussian filtering

Laplacian of Gaussian vs Derivative of Gaussian



Laplacian of Gaussian filtering



Derivative of Gaussian filtering

Zero crossings localize edges more accurately but are less convenient to use.

A Computational Approach to Edge Detection

JOHN CANNY, MEMBER, IEEE

Abstract—This paper describes a computational approach to edge detection. The success of the approach depends on the definition of a *comprehensive* set of goals for the computation of edge points. These goals must be precise enough to delimit the desired behavior of the detector while making minimal assumptions about the form of the solution. We define detection and localization criteria for a class of edges,

detector as input to a program which could isolate simple geometric solids. More recently the model-based vision system ACRONYM [3] used an edge detector as the front end to a sophisticated recognition program. Shape from motion [29], [13] can be used to infer the structure of


Canny Edge Detector

Adds non-maximum suppression and hysteresis thresholding to find edges from gradient outputs.

Gradients to Edges



original*
image



gradient
magnitude



gradient magnitude
after thresholding




How to turn these
thick regions of the
gradient into curves?



edge image

* Smoothing to suppress noise. Increase contrast to enhance edges.

Edge Thinning via Non-Maximum Suppression



For each pixel, check if it is a local max along the gradient direction. If so, keep as edge, if not, discard (set to 0).



Note that this requires check interpolated values at pixel locations p and r !

Or simply approximate by finding the closest pixel locations to p and r .

Non-Maximum Suppression Example

Magnitude


2	5	5	6	6	7	5	7
6	5	4	18	25	26	17	10
1	3	31	124	198	231	219	198
2	9	78	218	239	252	264	270
7	14	104	244	160	36	53	72
6	14	119	244	142	13	10	13
6	16	135	242	124	6	6	10
7	18	149	242	110	2	9	4



No interpolation approximation

Orientation


90	79	79	121	90	45	79	124
108	63	90	131	97	67	80	90
-135	-18	138	128	103	86	86	86
-180	117	151	148	121	88	85	85
135	146	169	172	168	96	79	81
162	172	176	177	179	-135	-90	-72
-180	173	175	175	175	-180	-59	-53
146	167	176	175	173	90	-27	-45



$$x = \tan\left(\frac{31}{180} \cdot \pi\right); x \approx 0.6$$

$$0.6 \cdot (124 - 198) + 198 \approx 154$$

$$0.6 \cdot (36 - 160) + 160 \approx 86$$



Keep as part of edge

Discontinuous Edges




gradient magnitude
after thresholding



edge image

Problem: pixels along
this edge didn't survive
the thresholding



Hysteresis thresholding: use a high threshold to start edge curves, and a low threshold to continue growing them.


Hysteresis Thresholding

White: Clears high threshold

Grey: clears low threshold but not high threshold

Mark pixels that clear high threshold as boundary


If there are pixels that clear low threshold and are connected to a boundary, mark them as boundary too



Hysteresis Thresholding




high threshold
(strong edges)



low threshold
(weak edges)

03. Gradients & Edges




hysteresis threshold


Canny Edge Detector

1. Filter image with derivative of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression (on low and high threshold):
 - Thin wide “ridges” down to single pixel width
4. Linking and thresholding (hysteresis):
 - Use the high threshold to start edge curves and the low threshold to continue them


Canny Step-by-Step



1: Filter image with derivative of Gaussian




gradients in y-dir




gradients in x-dir

2. Compute gradient magnitude and orientation




gradient magnitude




(visualization of) gradient
orientation


Canny Step-by-Step



3. Thresholding



gradient magnitude



(low) thresholded magnitude


4. Non-Maximum Suppression & Linking



Final Canny Edge Image



high threshold
(keeps only most confident pixels)



low threshold
(produces spurious edges)

Summary

- Image edges are a compact way to convey information;
- Edges found by applying derivative filter (e.g. Sobel) to extract gradients
- Gradients have two components: magnitude & orientation
- Smoothing prior to estimating gradients reduces effects of noise, but trades off localization accuracy
- Canny edge detection converts gradients to edges w/ non-max suppression & hysteresis thresholding