

Readings

- [Klette](#) 5, 5.2 Intro to segmentation, mean-Shift Segmentation
- [Szeliski](#) 5.3, 5.3.1, 5.3.2 k-Means, Mean-shift
- Forsyth & Ponce 9.3 Mean Shift Segmentation (uploaded online to LumiNUS)
- <http://www.chioka.in/meanshift-algorithm-for-the-rest-of-us-python/>

Summary

Motivation

- Segmentation separates image into coherent regions, sometimes to increase efficiency for further processing e.g. of recognition, other algorithms
- One interpretation of image segmentation is via clustering; each cluster is a segment
- To cluster, we need to consider how to:
 - Represent the points (feature representation)
 - Measure their differences (distance measure)
 - Do the clustering (clustering algorithm)

k-Means Clustering

- Iteratively assigns cluster membership and computing cluster centers
- Pros: simple, fast to compute, guaranteed to converge (to some local minimum)
- Cons: setting k non-trivial, sensitive to initialization, outliers, finds only spherical clusters
- To perform segmentation, represent each pixel as a feature (e.g. greyscale intensity, colour, colour + x,y position) and apply k-means clustering in feature space

SLIC Superpixels

- Superpixels are a group of local pixels that share common characteristics, e.g. intensity
- SLIC uses k-means to find superpixels based on their similarity in intensity and location

Mean-Shift Clustering

- The mean shift algorithm is a mode-seeking algorithm that finds local density maxima
- For each data point, iteratively shift a window towards its local centroid until arrival at some mode (attraction basin)
- Derivation shows that mean-shift performs gradient ascent to find the local maximum on the data distribution via a kernel density estimate
- To perform segmentation, apply mean shift to each point in feature space; assign all data points in an attraction basin to a single cluster
- Parameterized by bandwidth or window size h ; clustering results vary depending on h
- Pros: no assumption on shape, variable number of modes, robust to outliers
- Cons: h hard to set, computationally expensive
- Speedups possible by reducing the number of points we run mean-shift over

FAQs

Q: On slide 38 of lecture 5, points within radius r of a particular mode will belong to that mode. Is r necessarily smaller than or equal to the window size to get good clustering?

A: No strict rules -- depends on your kernel and the problem setting.

Q: On slide 20, there is the equation for composite distance. How do you manipulate the equation on the left to get the equation on the right? I attach my attempt below.

$$D = \left[\left(\frac{d_c}{d_{cm}} \right)^2 + \left(\frac{d_s}{d_{sm}} \right)^2 \right]^{1/2} \Rightarrow D = \left[d_c^2 + \left(\frac{d_s}{s} \right)^2 c^2 \right]^{1/2}$$

$$\begin{aligned}
 D &= \left[\left(\frac{d_c}{d_{cm}} \right)^2 + \left(\frac{d_s}{d_{sm}} \right)^2 \right]^{1/2} \\
 &= \left[\left(\frac{d_c}{c} \right)^2 + \left(\frac{d_s}{s} \right)^2 \times 1 \right]^{1/2} \\
 &= \left[\frac{d_c^2}{c^2} + \frac{d_s^2}{s^2} \times \frac{c^2}{c^2} \right]^{1/2} \\
 &= \frac{1}{c} \left[d_c^2 + \frac{d_s^2}{s^2} c^2 \right]^{1/2}
 \end{aligned}$$

A: You are correct there is this extra $1/c$ factor. However, because it is set as a constant, it simply scales the entire distance by this amount so we usually ignore it.

*Q: On slide 22, why do we consider $2s * 2s$ neighbour instead of $s * s$? Since each centre of a superpixel is s pixels away from each other, isn't $s * s$ neighbours (i.e. $0.5s$ on top, bottom, left, right) be sufficient? Beyond that, the pixel is closer to other neighbouring centres?*

A: The $0.5s$ only works well if your superpixels are regular (squares). If the shapes become more irregular this gives more flexibility. Note that this $2s$ is simply a recommended value trading off between efficiency (i.e. less points given to the clustering algorithm) vs. rate of change over the iterations (limiting to the $0.5s$ range)

Q: How does changing the kernel shape impact the mean shift algorithm?

A: Consider how the output will differ if you use a square kernel vs. a circular one.

Q: I don't get why we are required to formulate it as such

Use the smallest eigenvalue as the response function

$$R = \min(\lambda_1, \lambda_2)$$

For us, $A = H$ is a 2×2 matrix; we can directly solve for the eigenvalues:

$$\lambda_{\pm} = \frac{1}{2} \left[(h_{11} + h_{22}) \pm \sqrt{4h_{12}h_{21} + (h_{11} - h_{22})^2} \right]$$

Wouldn't we always pick λ_- - since it is always $< \lambda_+$?

A: Yes, but eigenvalues are rarely associated with the explicit definition which is only for 2×2 matrices. There are other ways of solving for the eigenvalues. Additionally, the theory for the second moment matrix is given in terms of the smaller eigenvalue.