

# Local Features I:

# Keypoints

CS 4243 Computer Vision & Pattern Recognition

Angela Yao

# Recap & Outline

## Last Lecture

- Visual textures: property indicative of material and appearance
- Filter banks as feature representations
- Representing texture with textons and histograms of textons
- Perceived contours: texture plays a big role

## Today's Lecture

- Keypoints: locations in images for computing descriptors and matching
- Corners & Harris operator
- Equivariance & invariance
- Automatic scale selection
- Laplacian of Gaussian (LoG) operator

San Francisco - 35mm film panorama from year 2000  
<https://www.flickr.com/photos/parkstreetparrot/27691676929/>

# Before Panorama Apps on Smartphones...

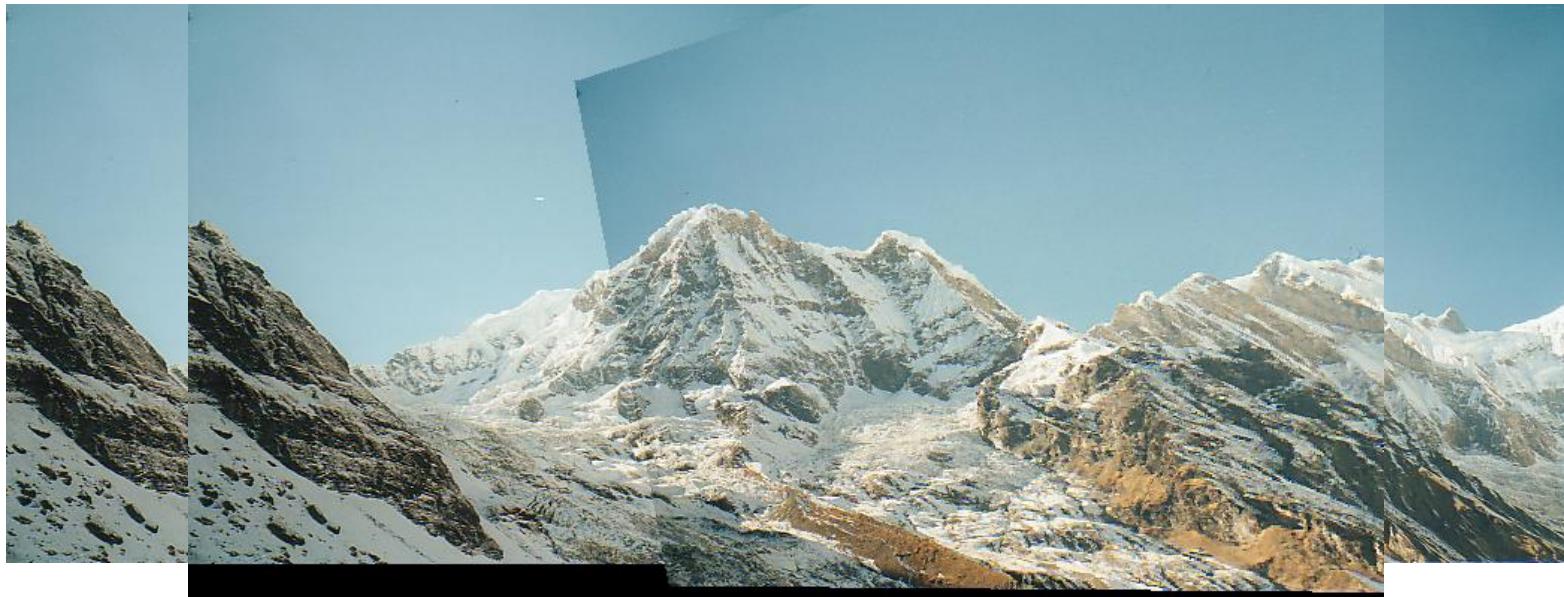


View from Lookout Mountain, Tennessee (1864)  
[https://en.wikipedia.org/wiki/Panoramic\\_photography](https://en.wikipedia.org/wiki/Panoramic_photography)



# How can we combine two images?

1. Match parts which are the same on both images
2. Align images based on the matches



# Matching is easy. (sort of)



by [Diva Sian](#)



by [swashford](#)



by [scgbt](#)

# What about these images?

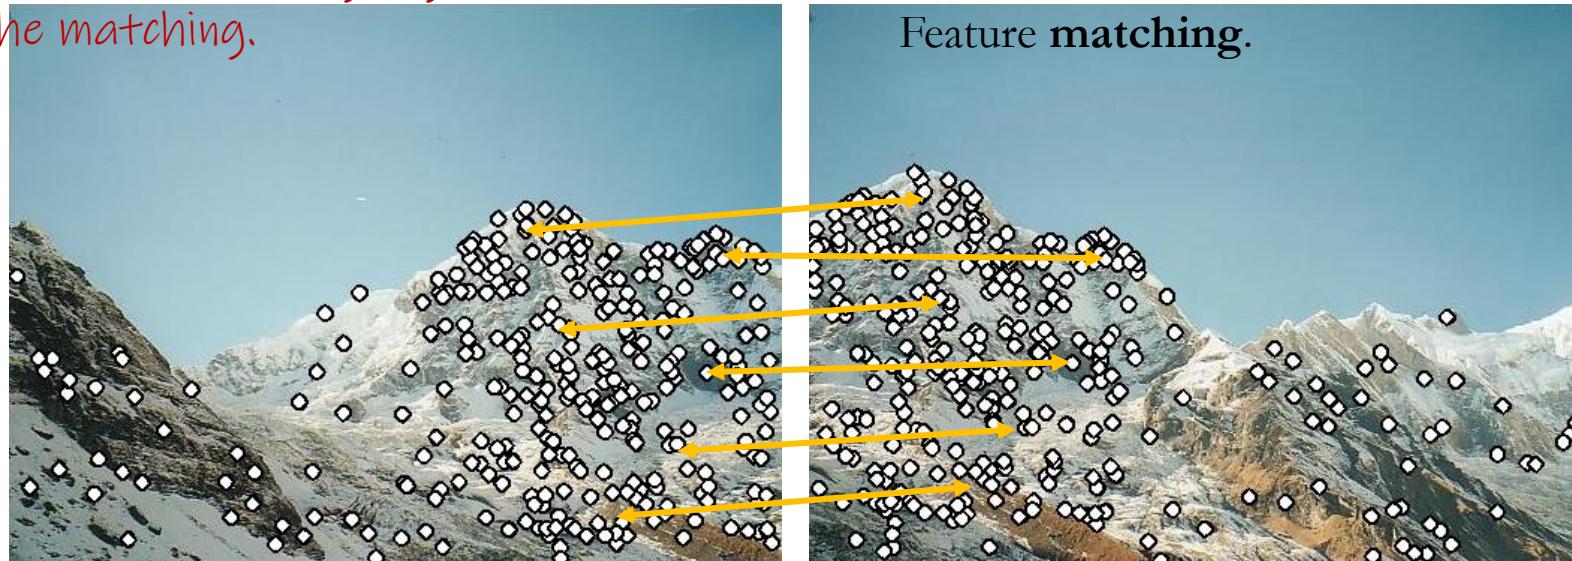


NASA Mars Rover images with SIFT feature matches

# How can we combine two images?

## 1. Match parts which are the same on both images

- a. Find locations to match.
- b. Represent surrounding region mathematically.
- c. Do the matching.



today's  
lecture

next  
week

Interest point **detection**.  
Compute feature **descriptors**.  
Feature **matching**.

Loaded terminology:  
aka keypoint detection; or  
simply feature detection

We will focus on a specific  
type of keypoint or interest  
point called Harris corners.

"Features"

...  
Often refers to  
both detection and  
representation, i.e.  
both the keypoint  
and the descriptor.

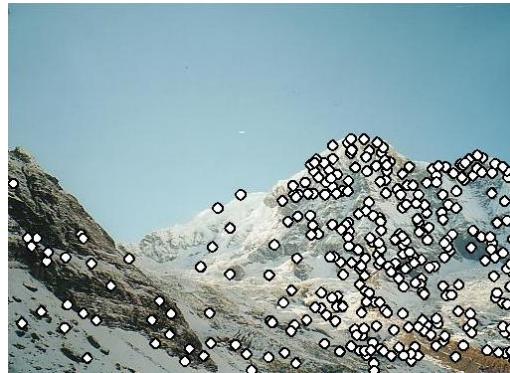
Isn't this easy? Just look for stuff that's the same right?  
Make a filter or something?

# Matching Two Images

# Matching Two Images

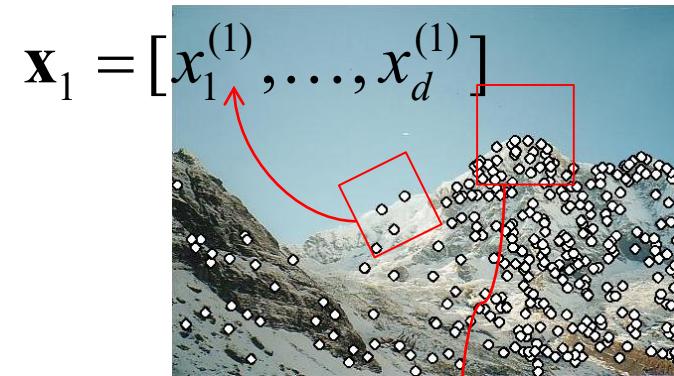
- 1) Detection: Identify interest points

Where should these come from? How to design these so that they will be found in both images?



- 2) Description: Compute vector feature descriptor surrounding each interest point.

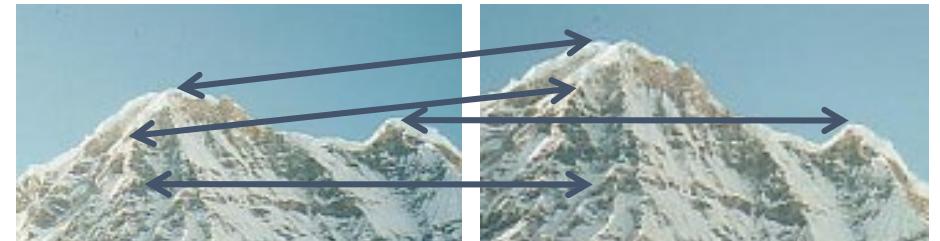
How should we characterize a region mathematically?  
How can we build in (enough but not too much)  
uniqueness to facilitate matching?



$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$

- 3) Matching: Determine correspondence between descriptors in two views

How can we establish correspondence? Efficiency?



9

# Characteristics of Good Local Features

- Repeatable interest points

The same point can be found in several images despite geometric and photometric transformations

- Distinct descriptors

Captures recognizably different information that allows it to be distinguished

- Efficient

Quantity: (number of features << number of image pixels); compact

- Local

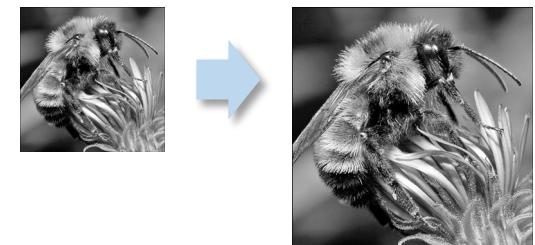
A feature occupies a relatively small area of the image (therefore robust to clutter and occlusion).

Geometric transformations

**Rotation**



**Scale**



Photometric transformations

**Intensity change**



# Goal: Repeatability of Interest Point Detector

We want to detect (at least some of) the same points in both images.

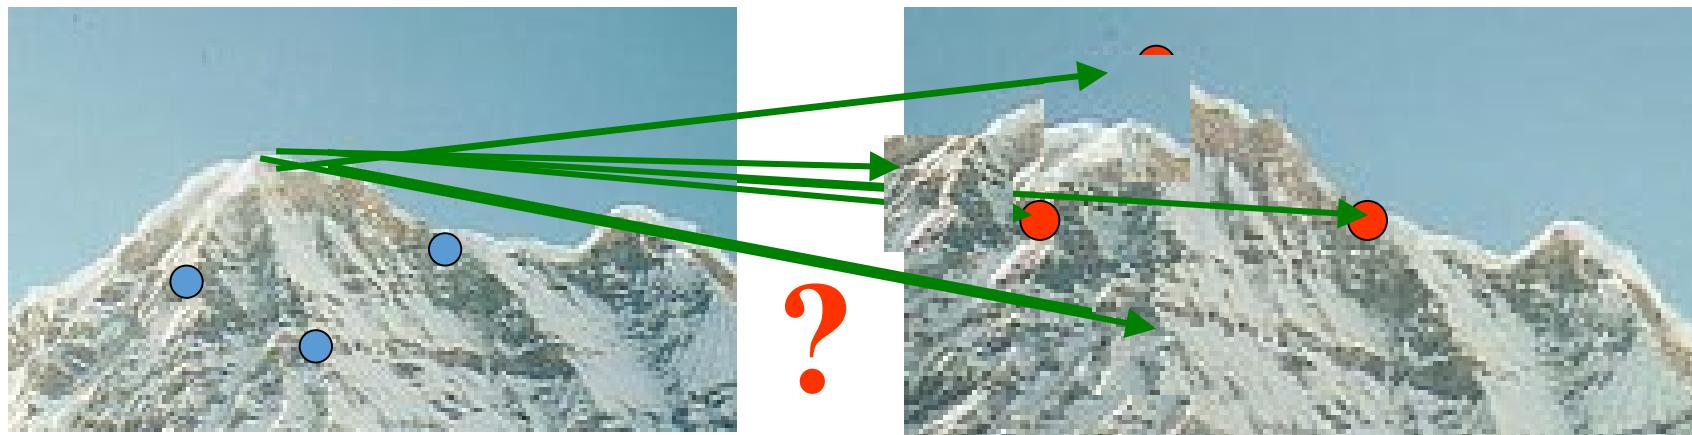


If the points are not the same, there is no chance to find true matches!

Yet we have to be able to run the detection procedure *independently* per image.

# Goal: Descriptor Distinctiveness

We want to reliably determine which point goes with which.



Descriptors must provide some invariance to geometric and photometric differences between the two views.

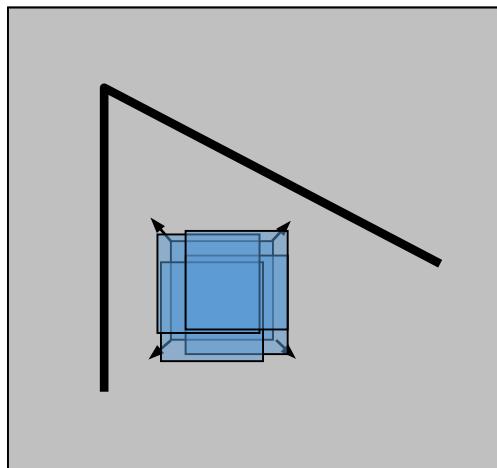
# Goal: Descriptor Distinctiveness

How to define?

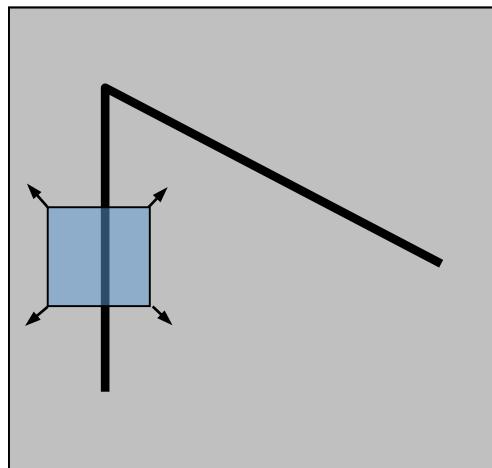
Look for unusual regions → leads to unambiguous matches in other images

Start by considering a small region: What makes a good candidate?

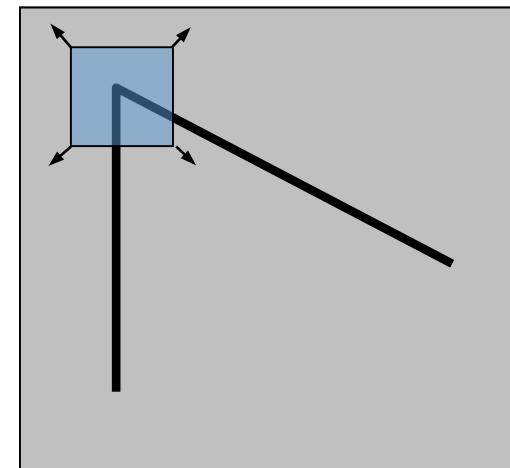
Define a window at the region. How will this window change when you shift it? At unique regions, shifting the window in any direction causes a big change.



“flat” region  
no change in all  
directions



“edge”  
no change along the  
edge direction



“corner”  
significant change in  
all directions

rationale and  
naming behind  
Harris corners

# The Math Behind Corner Detection

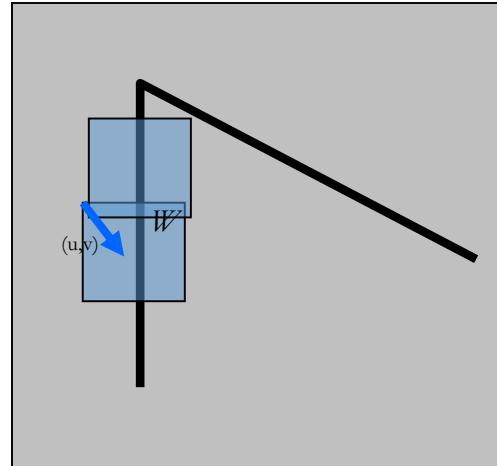
# The Math Behind Corner Detection

Consider shifting the window  $W$  by offset  $(u, v)$

- how do the pixels in  $W$  change?
- compare each pixel before and after by summing up the squared differences (SSD)
- this defines an SSD “error”  $E(u, v)$ :

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

- We are happy if this error is high. ←



Now just search for high error windows all over the image to find the corners.  
Efficiency issues?

HwQ: How does this scale computationally? Consider for a window of  $m \times m$  which we shift in by  $m$  offsets in both  $x$  and  $y$  for an image of  $w \times h$ .

# Assumption: consider only small shifts

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

Taylor Series expansion

$$I(x + u, y + v) = I(x, y) + \frac{\partial I}{\partial x} u + \frac{\partial I}{\partial y} v + \cancel{\text{higher order terms}}$$

For small shifts, i.e.  $(u, v)$  are small,  
then 1<sup>st</sup> order approx. is good enough.

notation  
convenience

$$\frac{\partial I}{\partial x} = I_x$$

$$\begin{aligned} &\approx I(x, y) + \underline{\frac{\partial I}{\partial x}} u + \underline{\frac{\partial I}{\partial y}} v \\ &= I(x, y) + I_x u + I_y v \end{aligned}$$

Image gradients  
in x- & y-direction!

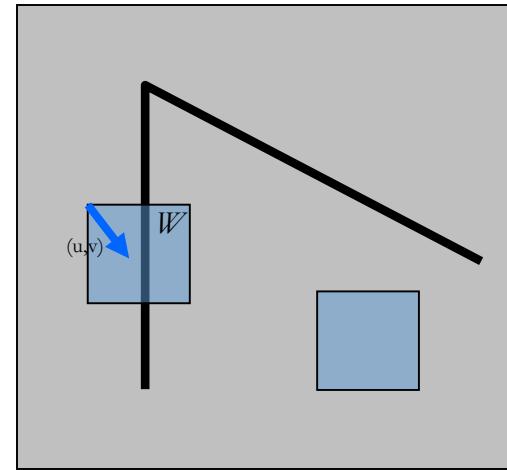
# The Math Behind Corner Detection

Consider shifting the window  $W$  by offset  $(u, v)$

- define an SSD “error”  $E(u, v)$ :

$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} [I(x, y) + I_x u + I_y v - \cancel{I(x, y)}]^2 \\ &= \sum_{(x,y) \in W} [I_x u + I_y v]^2 = Au^2 + 2Buv + Cv^2 \end{aligned}$$

$$= \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_{H \text{ Second moment matrix}} \begin{bmatrix} u \\ v \end{bmatrix} \quad A = \sum_{(x,y) \in W} I_x^2$$



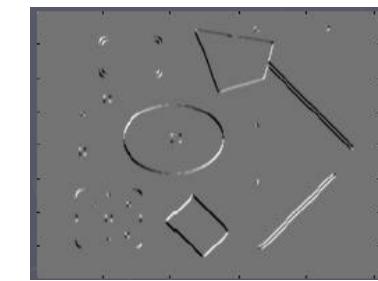
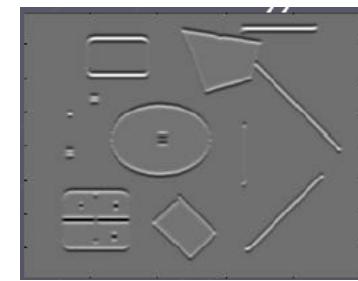
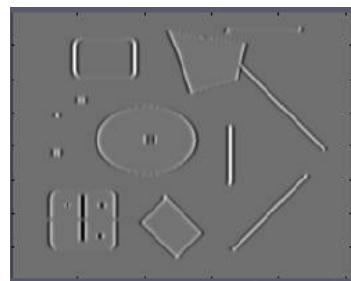
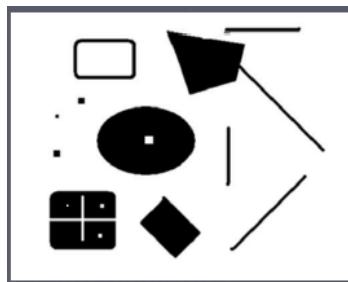
SSD error is locally approximated as a quadratic function!

$$B = \sum_{(x,y) \in W} I_x I_y \quad C = \sum_{(x,y) \in W} I_y^2$$

# Second Moment Matrix H

$$H = \begin{bmatrix} A & B \\ B & C \end{bmatrix} \quad A = \sum_{(x,y) \in W} I_x^2 \quad B = \sum_{(x,y) \in W} I_x I_y \quad C = \sum_{(x,y) \in W} I_y^2$$

2 x 2 matrix of image derivatives (averaged in neighborhood of a point).



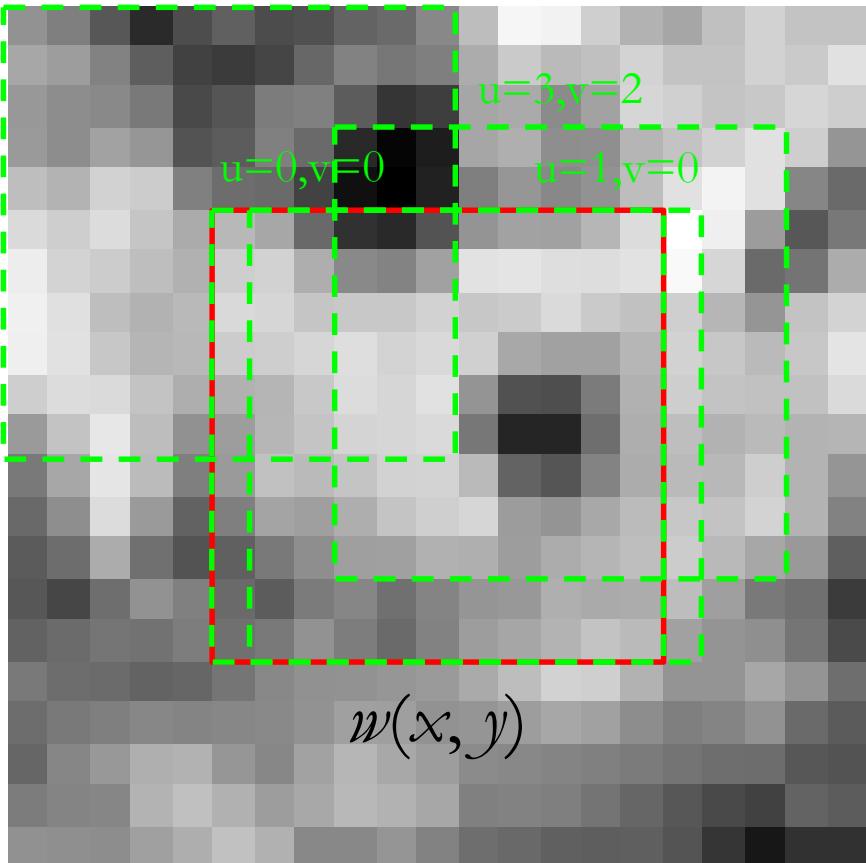
Notation:

$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$

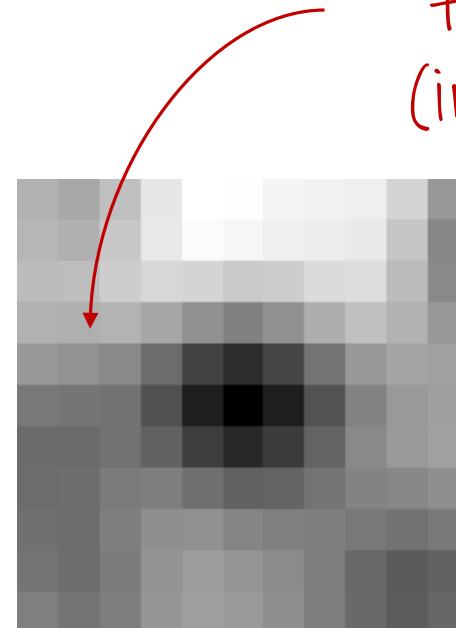
$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

$$I_x I_y \Leftrightarrow \frac{\partial I}{\partial x} \frac{\partial I}{\partial y}$$

# Visualizing the SSD Error



$$I(x, y)$$



$$E(u, v)$$

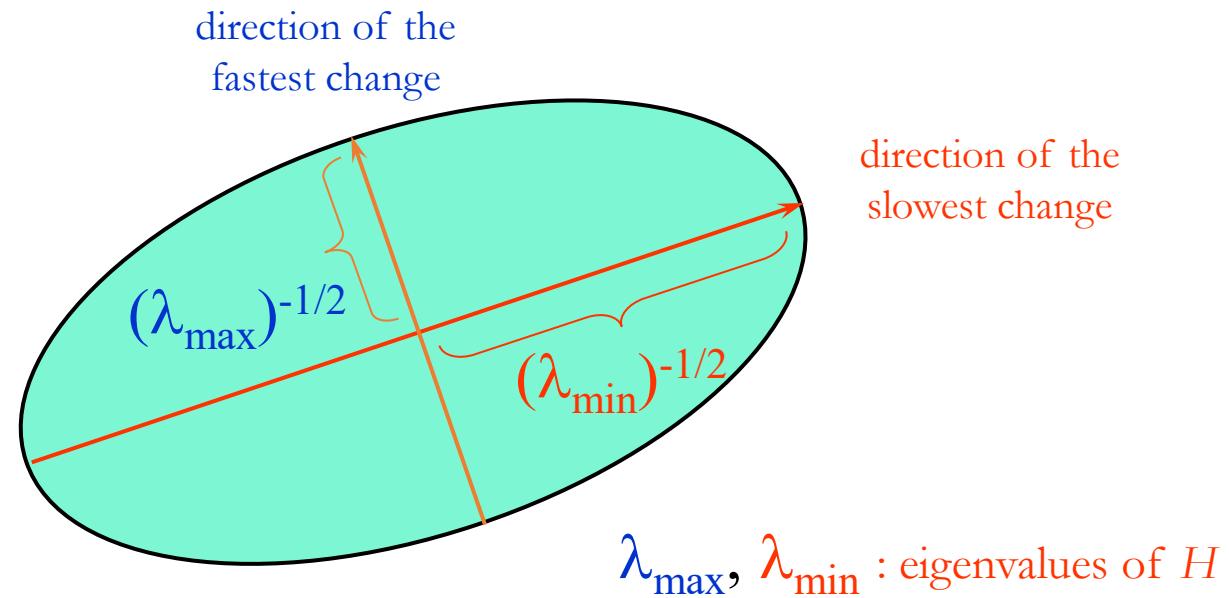
quadratic in  
the error  
(intensity).

# Visualizing H

$$[u \ v] H \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

Quadratic function can be visualized as an ellipse, where coefficients in  $H$  controls the shape of the ellipse.

Shape determined by looking at the eigenvectors and eigenvalues.



Ellipse axes orientation determined by the *eigenvectors* of  $H$

Ellipse axes length determined by the *eigenvalues* of  $H$

Solve with SVD (Singular Value Decomposition)!  
In  $2 \times 2$  case, directly.  
20

# Eigendecomposition of 2x2 Matrices

The **eigenvectors** of a matrix  $\mathbf{A}$  are the vectors  $\mathbf{x}$  that satisfy:

$$Ax = \lambda x$$

The scalar  $\lambda$  is the **eigenvalue** corresponding to  $\mathbf{x}$ .

$$\det(A - \lambda I) = 0$$

The eigenvalues are found by solving:

$$\det \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} = 0$$

For us,  $A = H$  is a 2x2 matrix; we can directly solve for the eigenvalues:

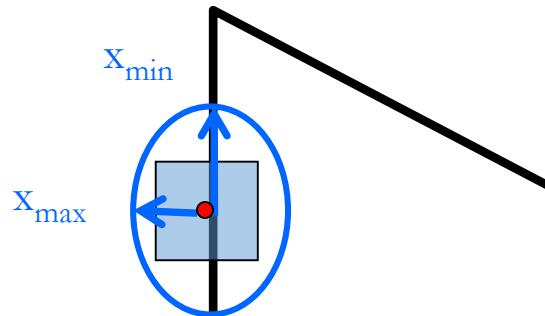
$$\lambda_{\pm} = \frac{1}{2} \left[ (h_{11} + h_{22}) \pm \sqrt{4h_{12}h_{21} + (h_{11} - h_{22})^2} \right]$$

Once you know  $\lambda$ , you find  $\mathbf{x}$  by solving

$$\begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = 0$$

# The Math Behind Corner Detection

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$



$$\begin{aligned} Hx_{\max} &= \lambda_{\max}x_{\max} \\ Hx_{\min} &= \lambda_{\min}x_{\min} \end{aligned}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$

Eigenvalues and eigenvectors of  $H$  define shift directions with the smallest and largest change in error

- $x_{\max}$  direction of largest increase in  $E$
- $\lambda_{\max}$  related to amount of increase in direction  $x_{\max}$
- $x_{\min}$  direction of smallest increase in  $E$
- $\lambda_{\min}$  related to amount of increase in direction  $x_{\min}$

If corners represent having large increase in all directions, this suggests that both  $\lambda_{\max}$  and  $\lambda_{\min}$  should be sufficiently large!

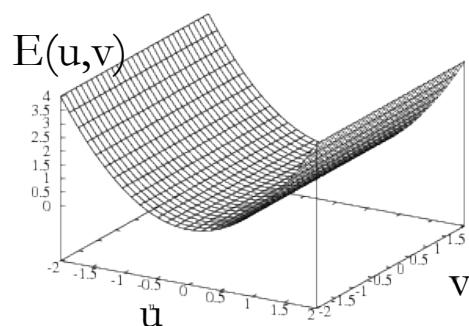
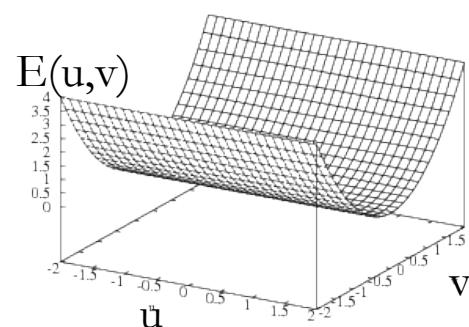
## Special Cases

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix}$$

$$H = \begin{bmatrix} A & B \\ B & C \end{bmatrix}$$

$$H = \begin{bmatrix} 0 & 0 \\ 0 & C \end{bmatrix}$$

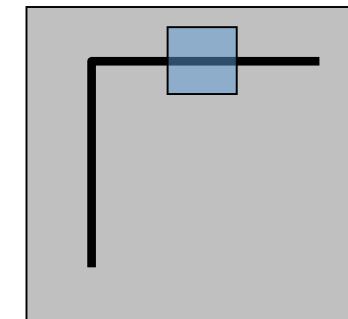
$$H = \begin{bmatrix} A & 0 \\ 0 & 0 \end{bmatrix}$$



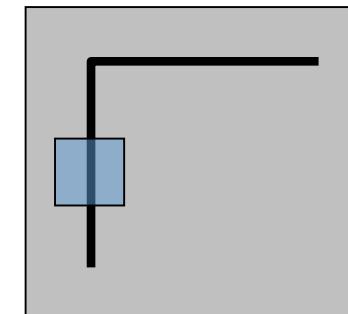
07. Local Features I: Keypoints

$$A = \sum_{(x,y) \in W} I_x^2 \quad B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$

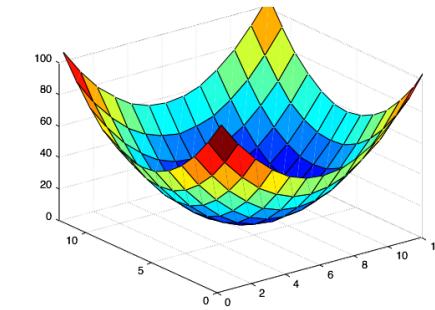
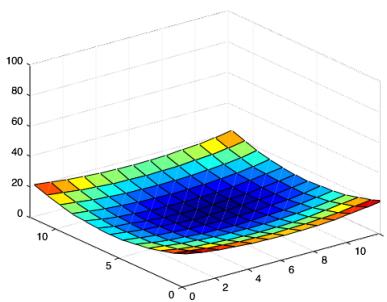
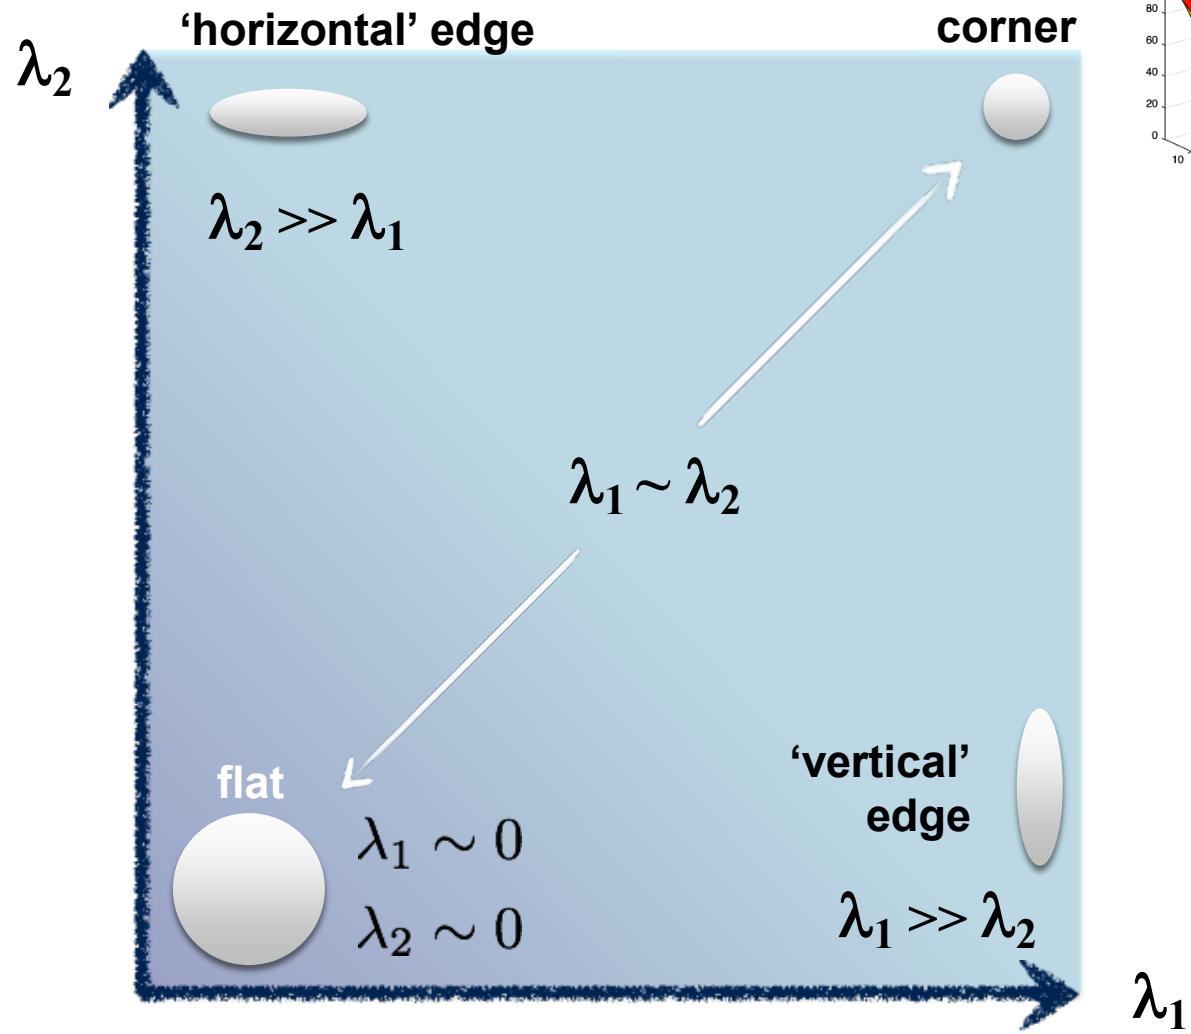
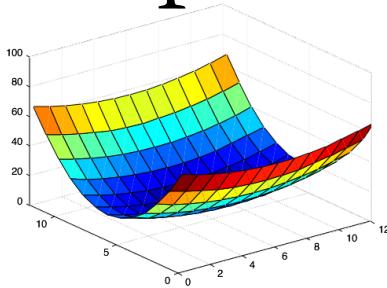


Horizontal edge:  $I_x = 0$

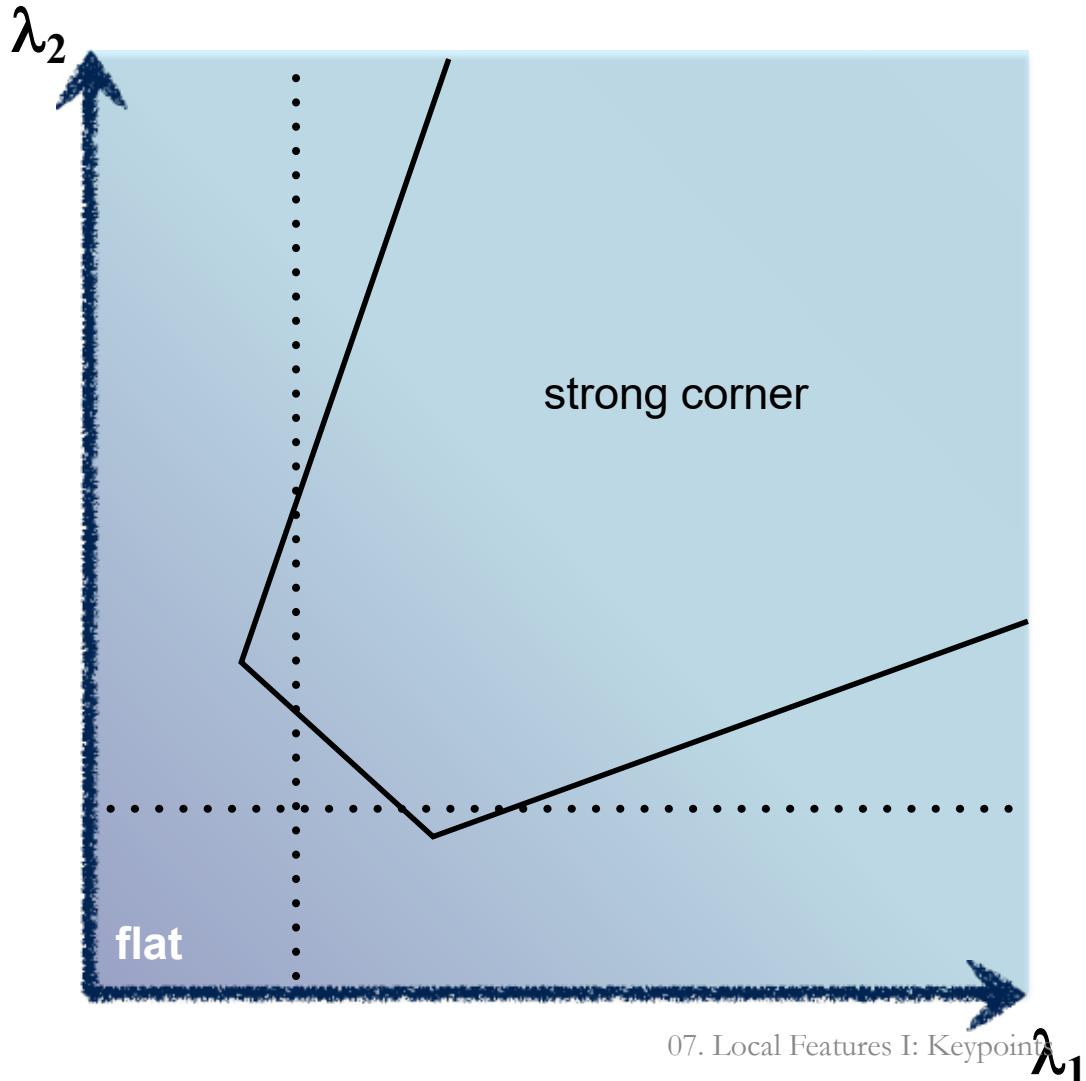


Vertical edge:  $I_y = 0$

# Interpreting the Eigenvalues of H



# Converting Eigenvalues to Corners



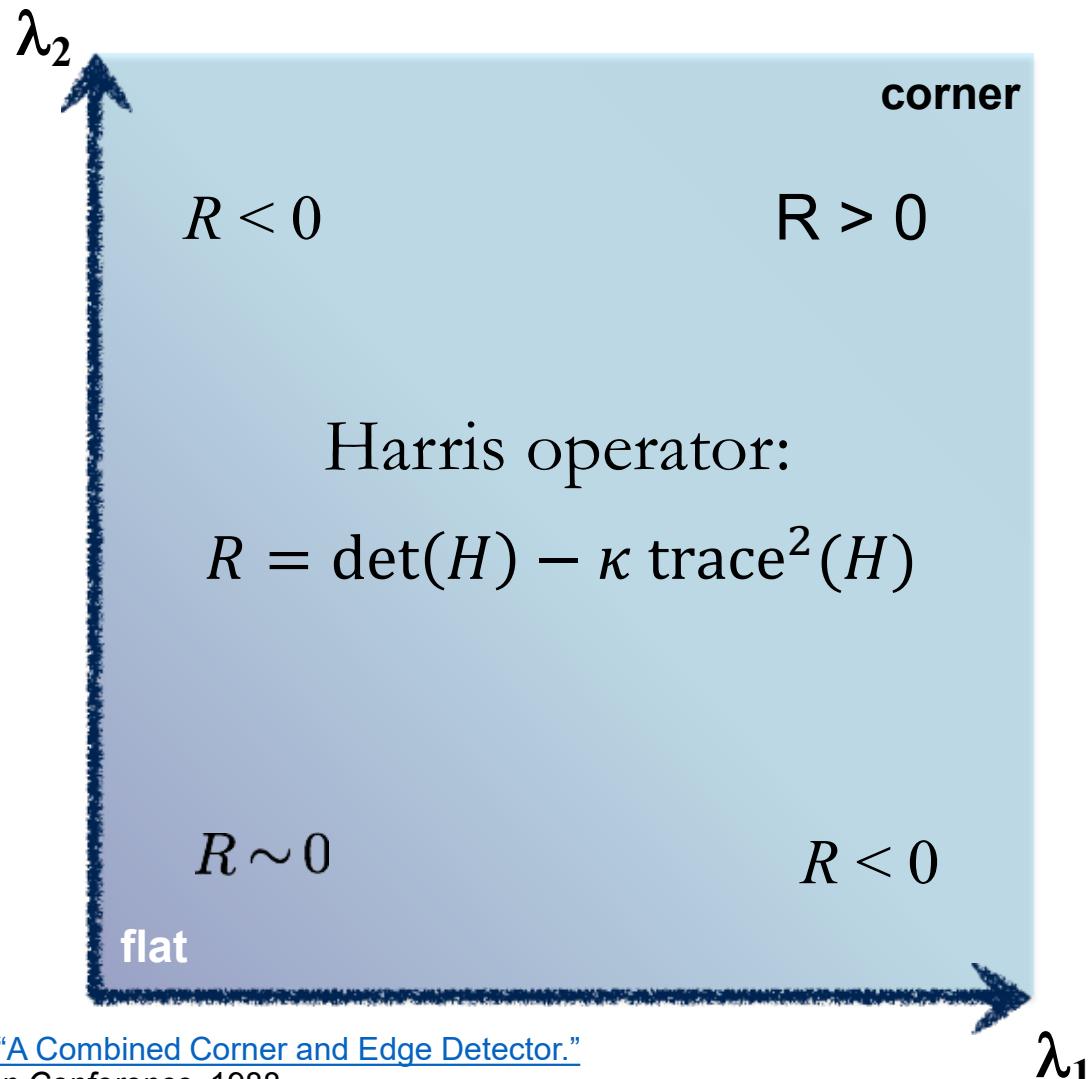
How can we develop a function to score ‘cornerness’?

Use the smallest eigenvalue as the response function

$$R = \min(\lambda_1, \lambda_2)$$

Getting eigenvalues is slow (even in 2x2 scenario, because we need to take square roots).

# Converting Eigenvalues to Corners (more efficiently)



A more efficient way is

$$R = \lambda_1 \lambda_2 - \kappa(\lambda_1 + \lambda_2)^2$$

$\kappa \sim 0.04 \text{ to } 0.06$

$$\det H = \lambda_1 \lambda_2$$

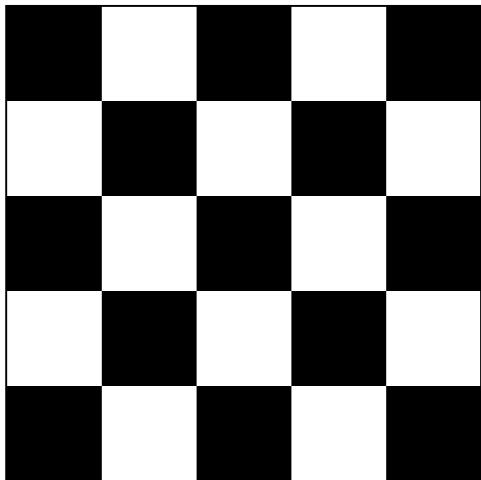
$$\operatorname{trace} H = \lambda_1 + \lambda_2$$

$$\det \begin{pmatrix} a & b \\ c & d \end{pmatrix} = ad - bc$$

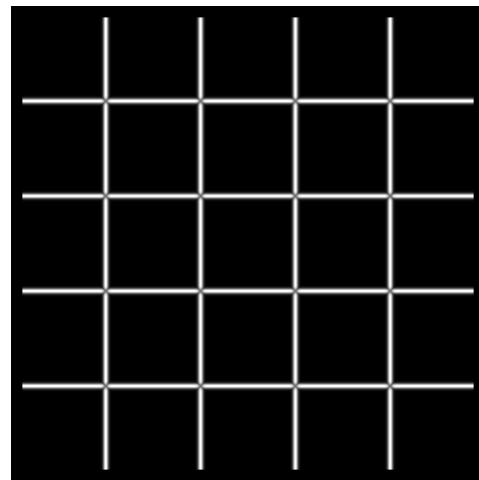
$$\operatorname{trace} \begin{pmatrix} a & b \\ c & d \end{pmatrix} = a + d$$

Eigenvalues are always greater than 0 by definition since the H matrix is positive semi-definite.

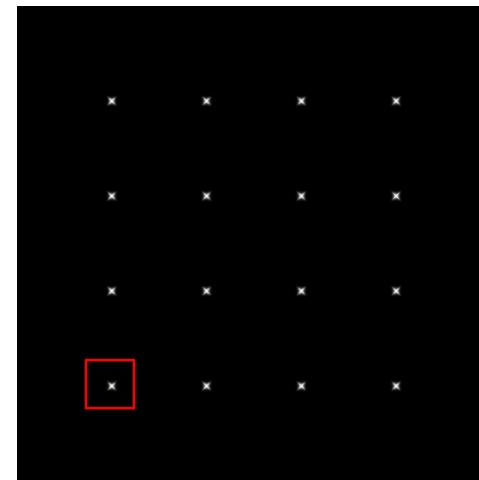
# Comparison of Corner Criteria



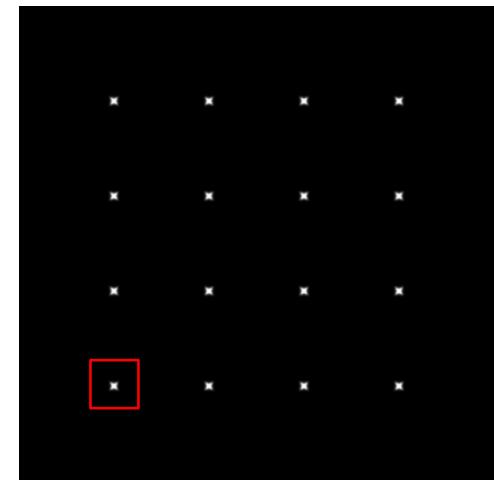
$I$



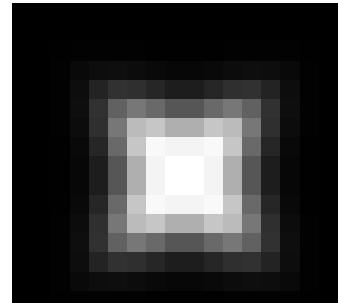
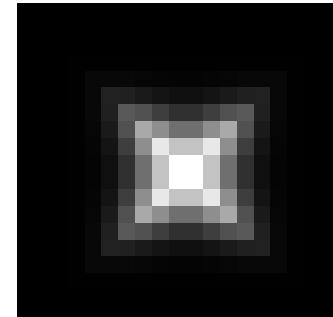
$\lambda_{\max}$



$\lambda_{\min}$



Harris operator

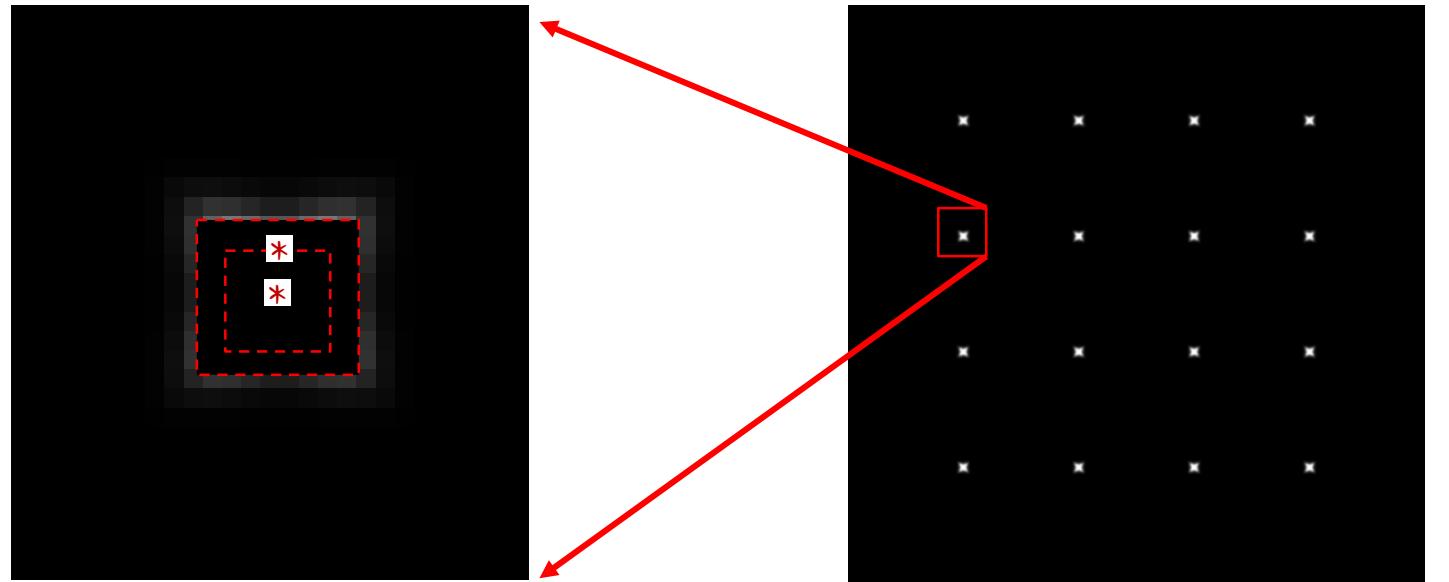


# Summary of Harris Corner Detector

- 1) Compute gradient at each point in the image
- 2) Compute  $H$  matrix for each image window and get their *cornerness* scores.
- 3) Find points whose surrounding window gave large corner response ( $R >$  threshold)
- 4) Take the points of local maxima, i.e., perform non-maximum suppression



# (Simple) Non-Maximum Suppression



Harris  
operator

*Speed-ups possible to make this more efficient.*

Simplest approach: iteratively search for “max” values,  
then zero-out everything in the surrounding window.

# Adaptive Non-Maximum Suppression

- Simple search of maxima leads to an uneven distribution of interest points in areas of higher contrast
- Pick corners which are both local maxima and whose response is significantly greater (e.g. 10%) than all neighbouring local maxima within some radius  $r$



(a) Strongest 250



(b) Strongest 500



(c) ANMS 250,  $r = 24$



(d) ANMS 500,  $r = 16$

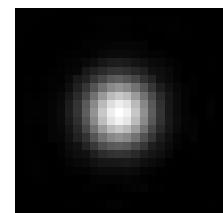
# Weighting the Derivatives

In practice, using a simple window  $W$  (e.g. 3x3, 5x5) doesn't work too well

$$H = \sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Instead, we *weight* derivatives based on its distance from the center pixel

$$H = \sum_{(x,y) \in W} w_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$



$$w_{x,y}$$

# Harris Corner Invariances

# Invariance and equivariance

A criterion of corner detection: repeatability.

Suppose the image is transformed in some way

- translation
- rotation
- intensity scaling

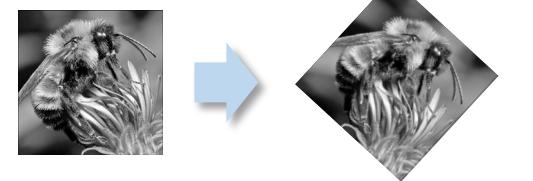
How *should* the corners change?

- Whether a corner is detected or not?
- Location?

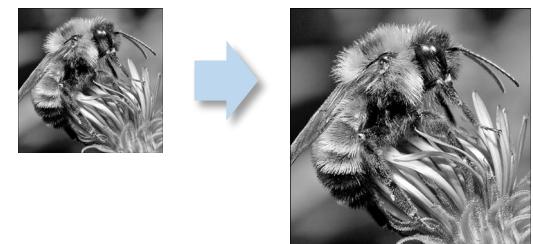
How *does* the output of Harris corner detector change?

Geometric transformations

**Rotation**

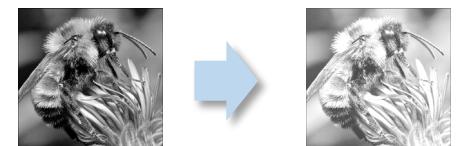


**Scale**



Photometric transformations

**Intensity change**



# Invariance and Equivariance

**Equivariance:** if we have two transformed versions of the same image, detection (locations) undergoes a similar transformation

**Invariance:** image is transformed and detection (score) does not change

“Detection” here can refer either to the corner location as well as the probability of being detected as a corner. Can think of equivariance wrt location and invariance wrt score.

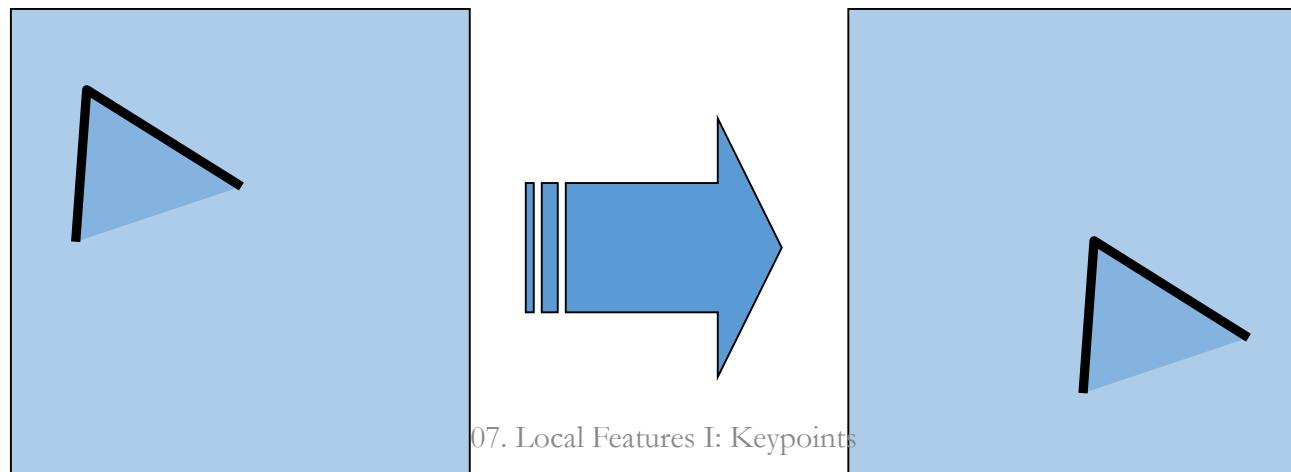
Sometimes invariance and equivariance are  
both referred to as invariance. ☹

# Invariant to translation?

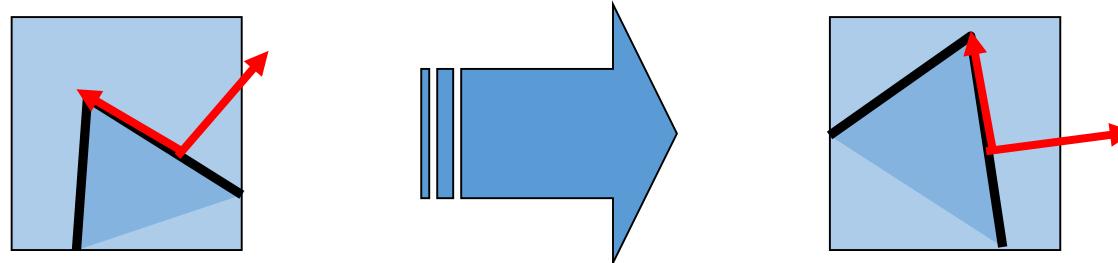
What happens if image is translated?

- Derivatives,  $H$  obtained through convolution,
- Eigenvalues based only on derivatives

Thus Harris corners; detected location is *equivariant to translation*, and response ( $R$ ) is *invariant to translation*



# Invariant to Rotation?



- Eigenvectors represent the direction of maximum / minimum change in appearance, so they rotate with the patch
- Eigenvalues represent the corresponding magnitude of maximum/minimum change so they stay constant

Corner response is only dependent on the eigenvalues → invariant to rotation

Corner location is as before equivariant to rotation.

# Invariant to Photometric Transformations?

Assume linear change:  $I' = aI + b$

Change to image gradients?

- $I'_x = aI_x$        $I'_y = aI_y$

Second Moment matrix?

- $H' = a^2H$

Eigenvalues & cornerness response function?

- $\lambda'_i = a^2\lambda_i$        $R' = a^2R$       *a<sup>2</sup> if R is min(λ<sub>1</sub>, λ<sub>2</sub>), a<sup>4</sup> if R is Harris operator*

Does this change probability of detecting a corner?

*Yes, since R' is now scaled by a<sup>2</sup>*

What happens if no scaling (i.e.,  $a = 1$ )?

*R' is not affected.*



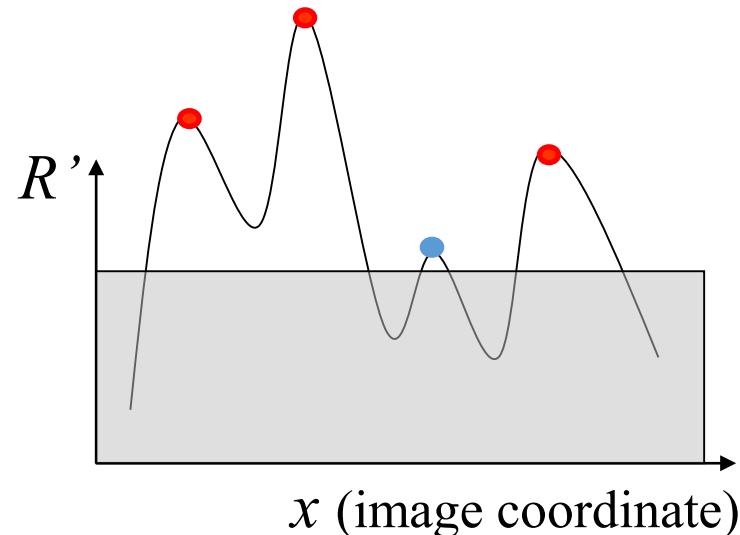
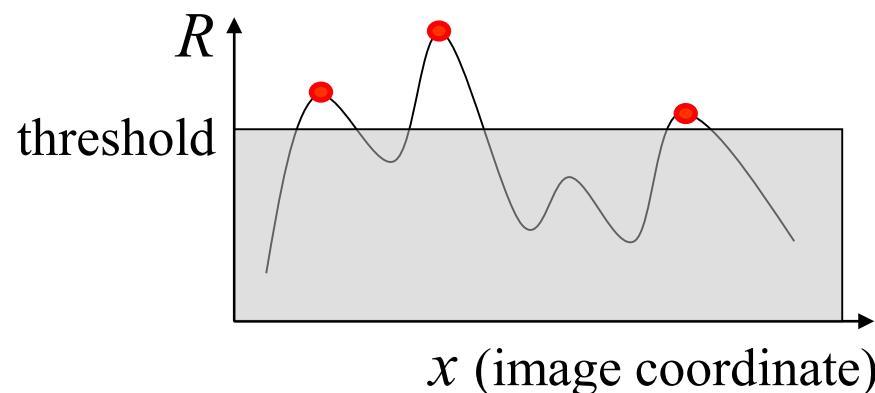
# Invariance to photometric transformations?

The Harris corner detector (both locations and probability of detection) is invariant to additive changes in intensity

Changes in overall “Brightness”

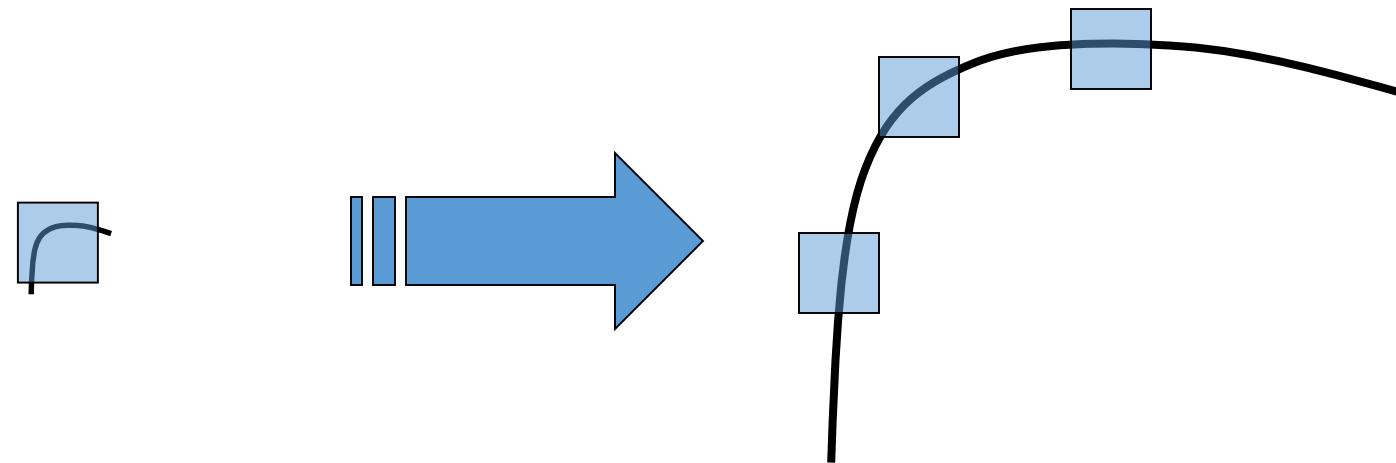
It is not invariant to scaling of intensity  $I' \rightarrow \alpha I$

Changes in “Contrast”



# Invariance to Scaling?

What was one patch earlier is now many.



Corner

all points will  
be classified as **edges**

Not equivariant to scaling

# Automatic scale selection

One definition of  $f$  could be the Harris operator ( $R$ ) ... another is the Laplacian of Gaussians (coming up)

When looking for keypoints, find the scale that gives a local maximum of  $f$

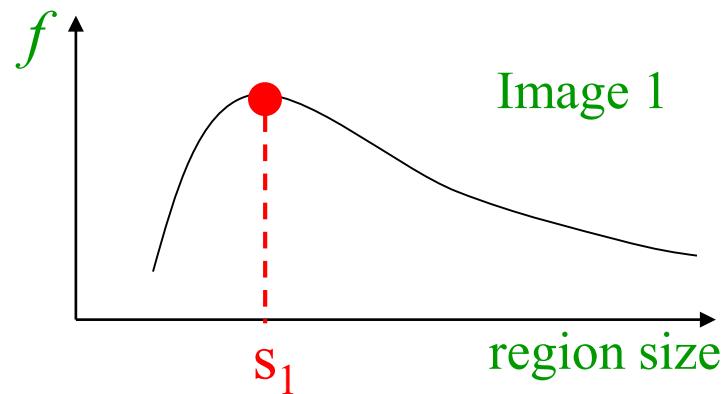
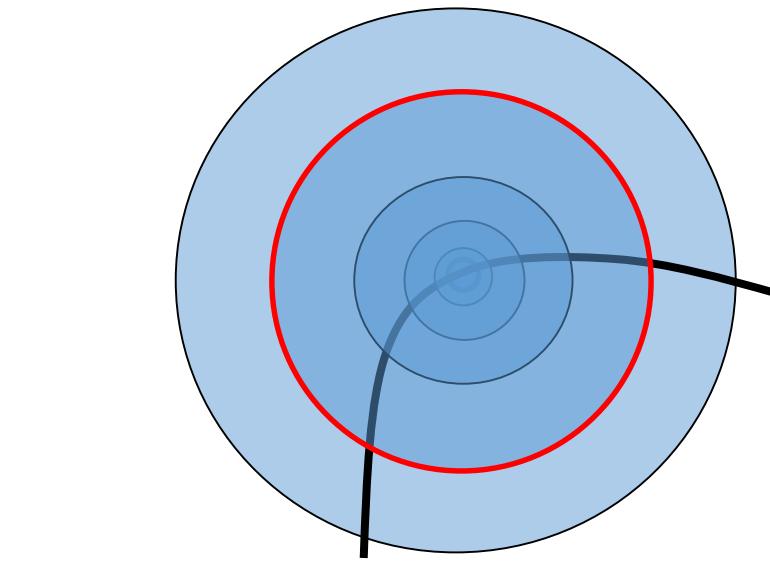


Image 1

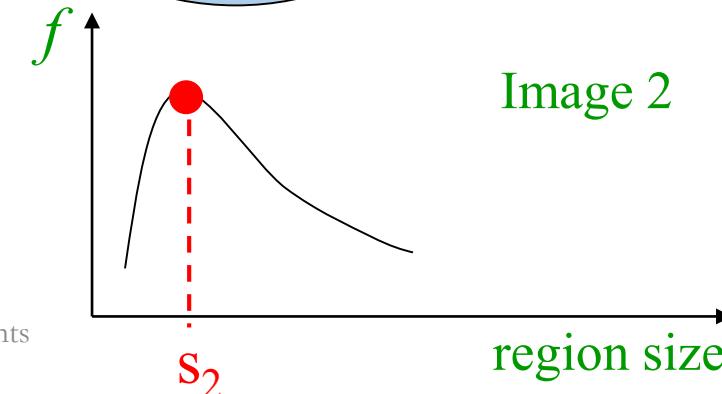
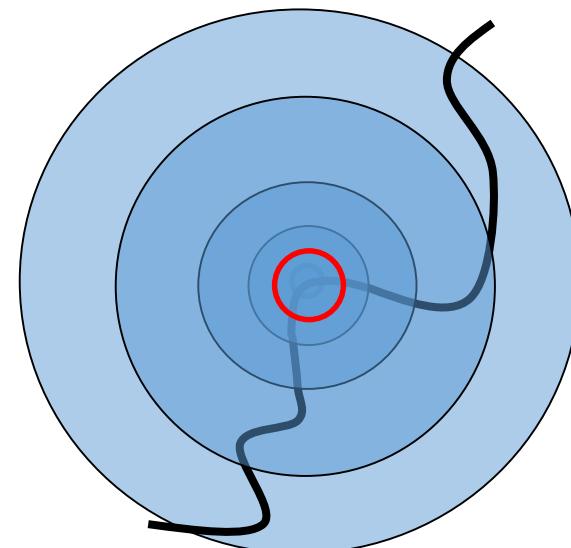
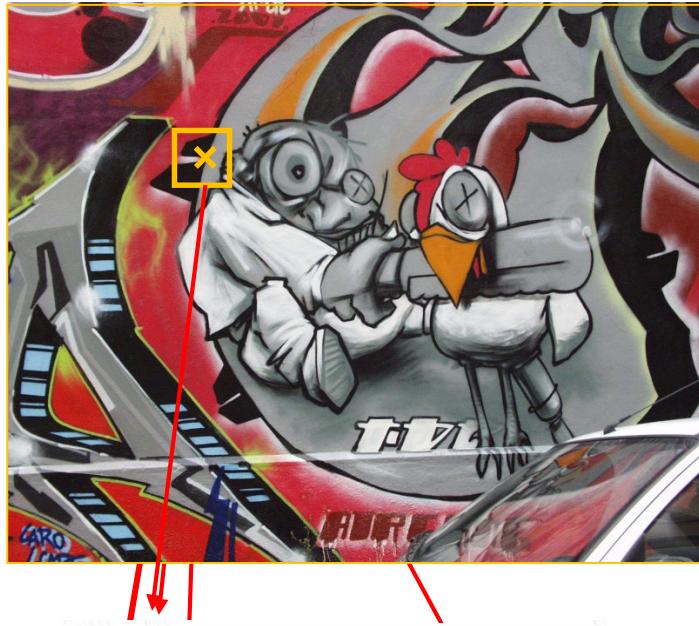
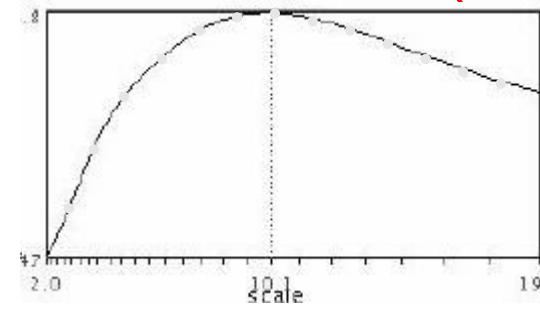
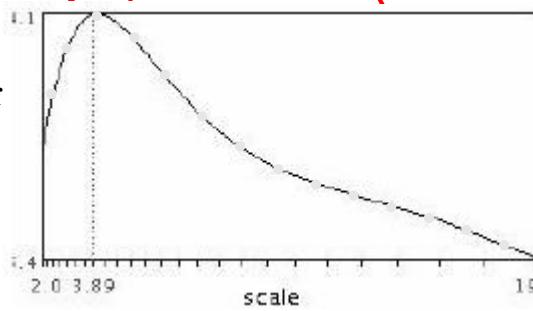


Image 2

# Automatic Scale Selection



Function  
responses for  
increasing  
scale (scale  
signature)



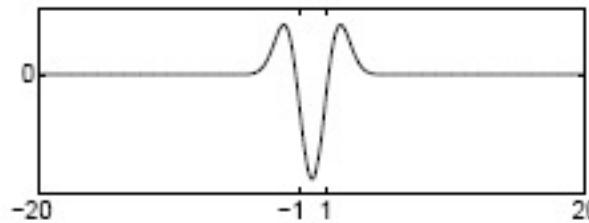
$$f(I_{i_1 \dots i_m}(x, \sigma)) = f(I_{i_1 \dots i_m}(x', \sigma'))$$

Same operator responses if  
the patch contains the same  
image up to some scale factor.

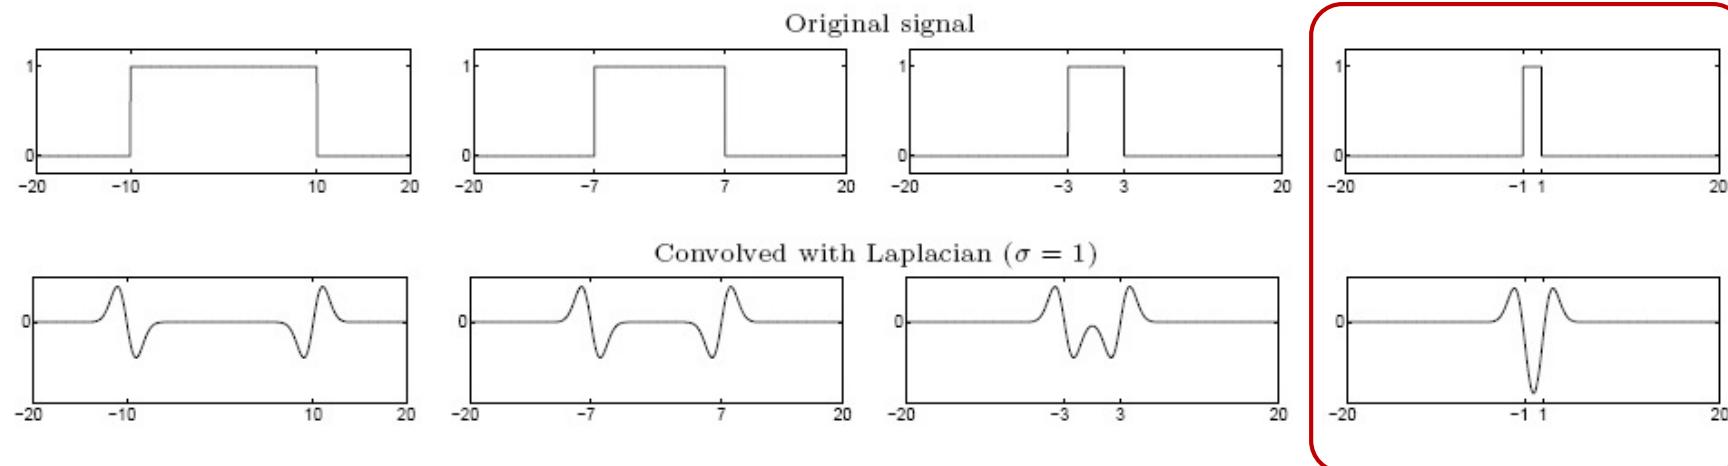
# LoG: Laplacian of Gaussians

“Blob detector”

# Laplacian of Gaussian (LoG) – 1D



2<sup>nd</sup> derivative of Gaussian

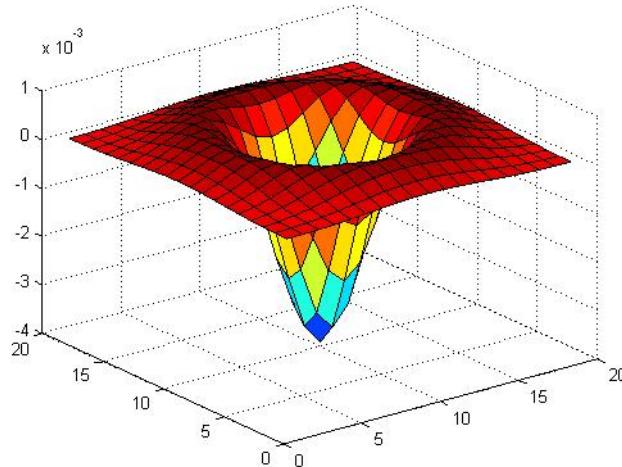


Highest response when the signal has the same  
**characteristic scale** as the Gaussian.

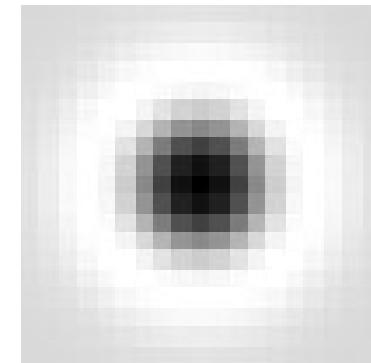
# Laplacian of Gaussian (LoG) – 2D

$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

Laplacian operator

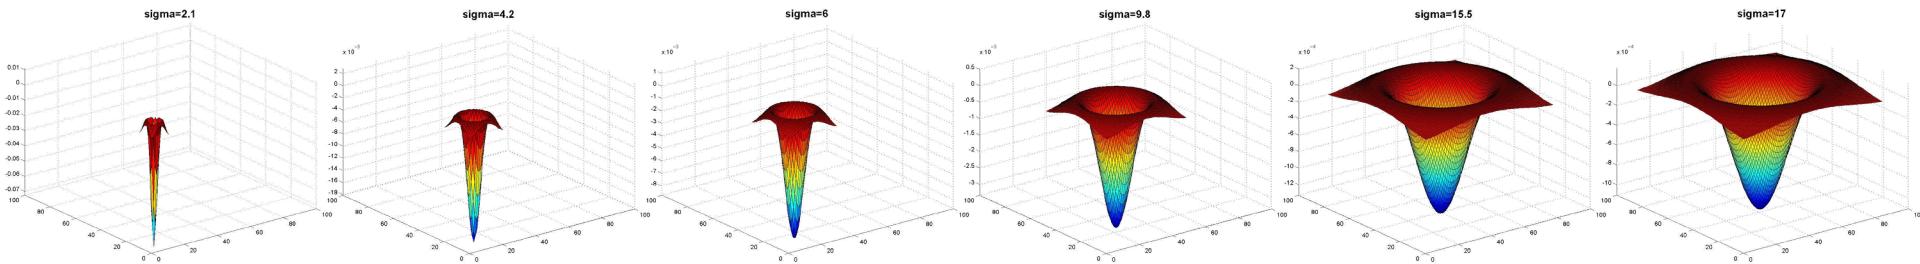


“Blob detector”



very similar to a difference of Gaussians – i.e. a Gaussian minus a slightly smaller Gaussian  
(cheaper approximation than taking 2<sup>nd</sup> derivative)

# Applying different Laplacian filters



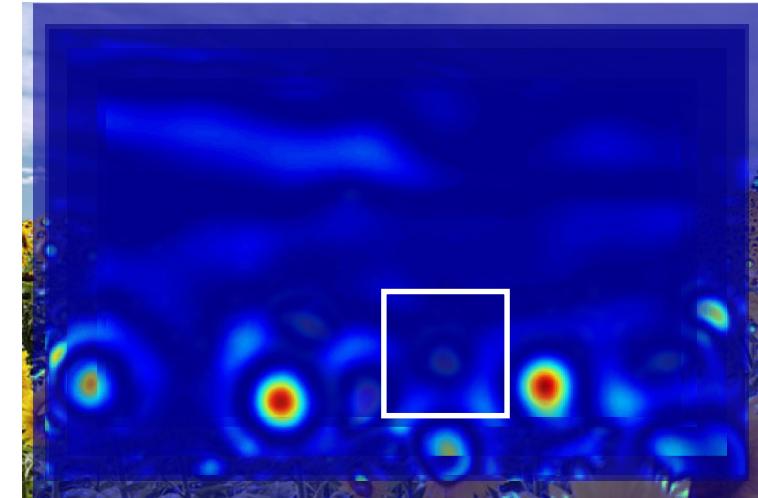
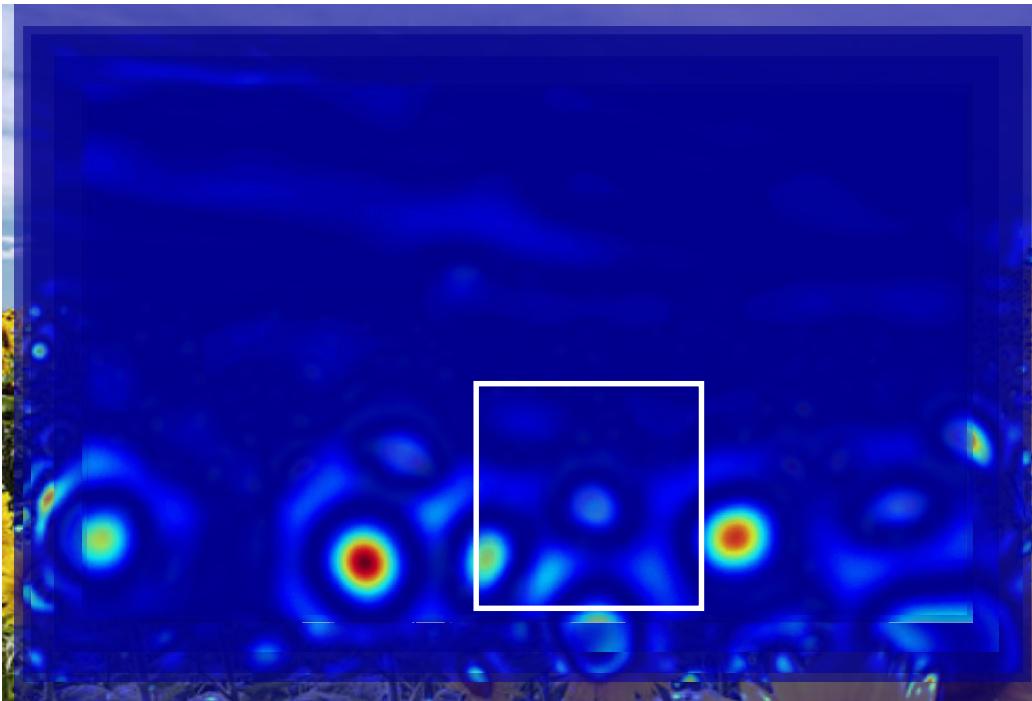
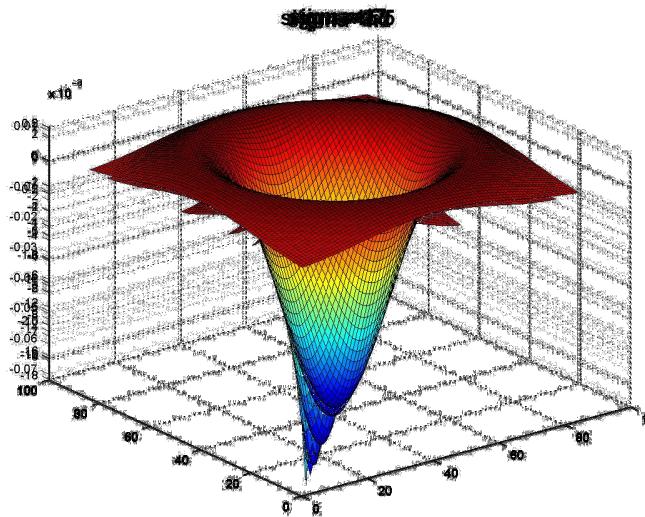
Full size



3/4 size



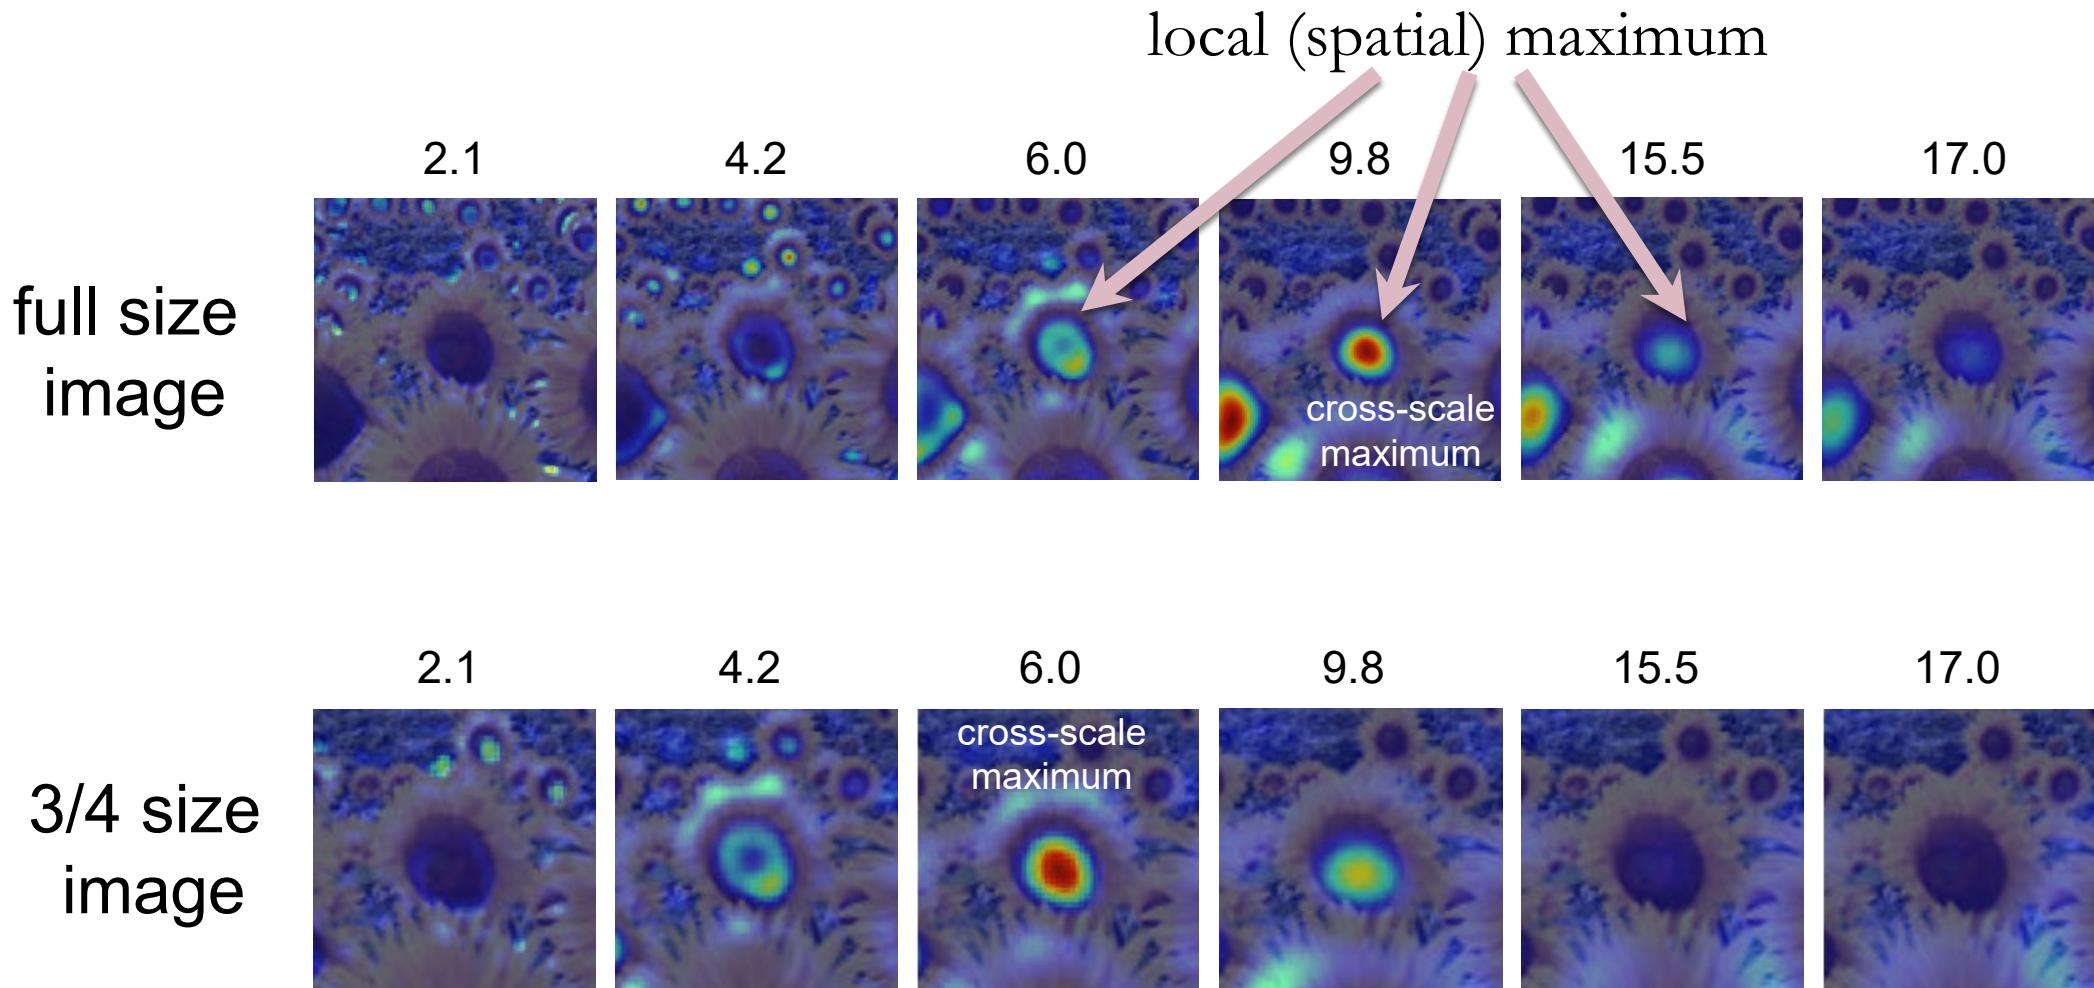
Focus on what happens to these 2 flowers.



**jet color scale**  
blue: low, red: high

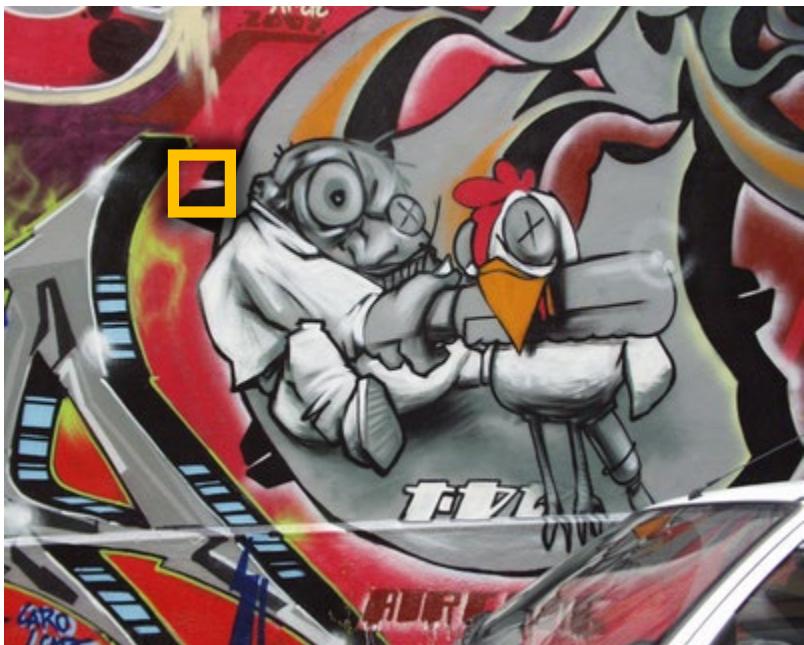
Adapted from K. Grauman, K. Kitani

# Finding the Optimal Scale

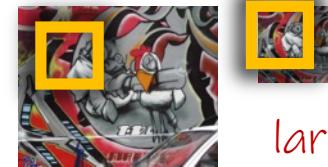


# Implementation

Instead of computing  $f()$  for larger and larger windows, we can implement using a fixed window size with a Gaussian pyramid



smallest scale



largest scale

# Summary

- “Corners” are good candidates as keypoints; can be found by looking at the Hessian matrix (eigenvalues are both large and approximately equal)
- Harris operator is one (relatively efficient) way of finding these corners
- Harris corners are equivariant to geometric transformations such as translation, rotation, semi-invariant to photometric transformations (brightness changes but not contrast) and not invariant to scaling
- To introduce robustness to changes in scale, we can use automatic scale selection as a part of finding keypoints
- The Laplacian of Gaussian (LoG) operator is an alternative keypoint detector which detects “blobs” and is scale-sensitive